bbob constrained f44 d10 i1 — what the agent chose, and when

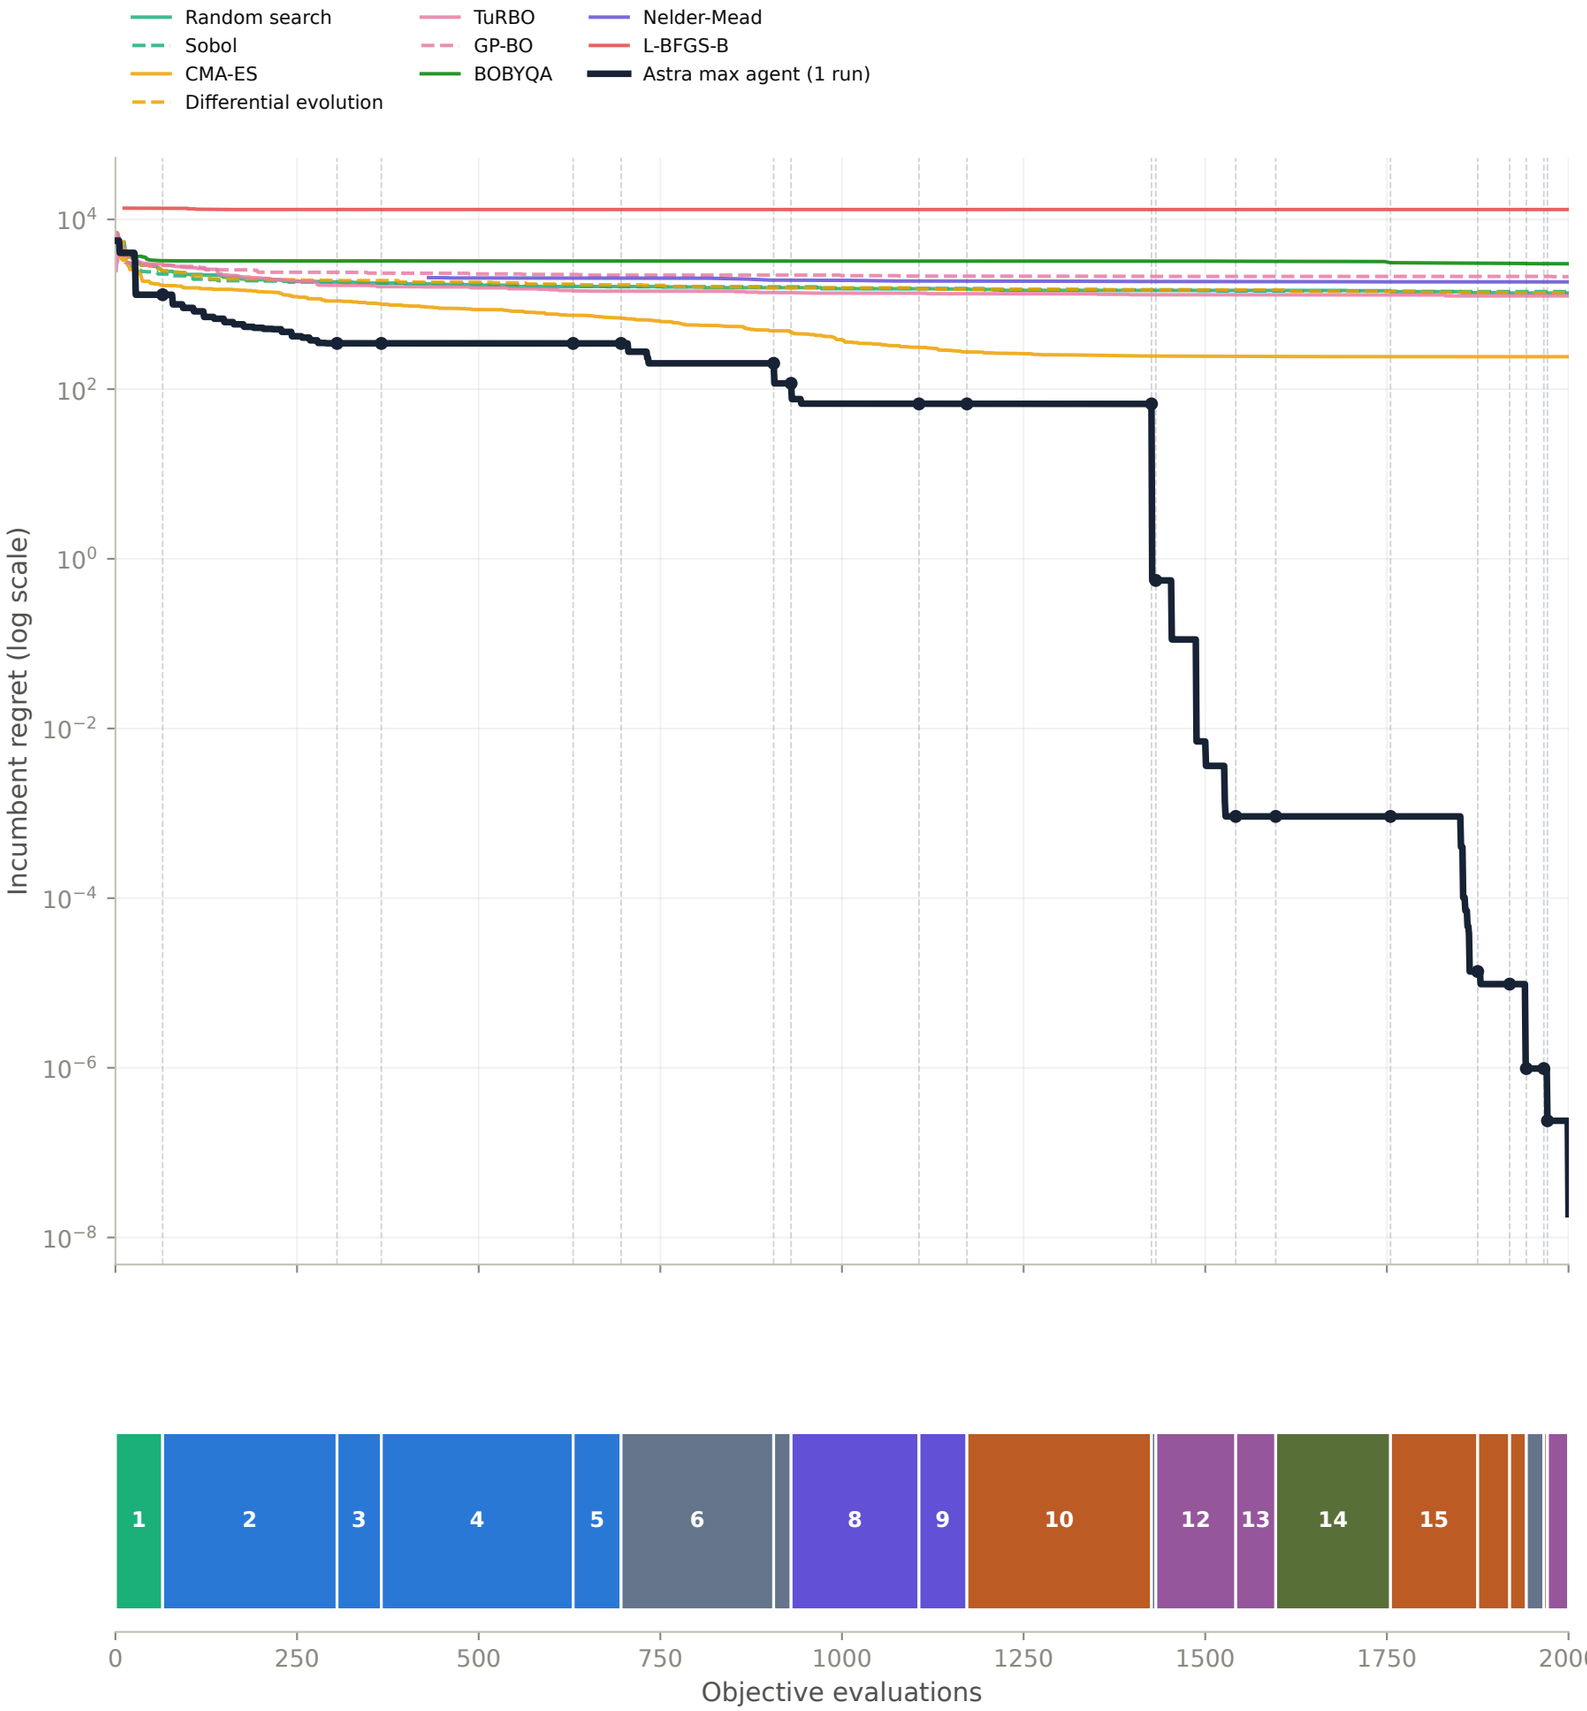

Block decisions

1. sobol_feasibility
   65 calls · evaluations 1–65
2. slsqp_local
   240 calls · evaluations 66–305
3. slsqp_local
   61 calls · evaluations 306–366
4. slsqp_restart
   264 calls · evaluations 367–630
5. slsqp_restart
   66 calls · evaluations 631–696
6. coordinate_grid
   210 calls · evaluations 697–906
7. additive_recombine
   24 calls · evaluations 907–930
8. surrogate_slsqp
   176 calls · evaluations 931–1106
9. surrogate_slsqp
   66 calls · evaluations 1107–1172
10. additive_slice_refine
    254 calls · evaluations 1173–1426
11. refined_recombine
    6 calls · evaluations 1427–1432
12. scaled_slsqp
    110 calls · evaluations 1433–1542
13. scaled_slsqp
    55 calls · evaluations 1543–1597
14. additive_extra_intervals
    158 calls · evaluations 1598–1755
15. cobyqa_boundary
    120 calls · evaluations 1756–1875
16. cobyqa_boundary
    44 calls · evaluations 1876–1919
17. boundary_projection
    23 calls · evaluations 1920–1942
18. tangent_newton_polish
    24 calls · evaluations 1943–1966
19. boundary_residual_repair
    5 calls · evaluations 1967–1971
20. fine_tangent_polish
    29 calls · evaluations 1972–2000

Boundaries: evaluator snapshots of prospective decisions; method changes inside a block stay in that block.
Agent: one fresh run, 2,000/2,000 evaluations; GPT-6 Astra, max reasoning. Baselines: seed means (10 stochastic / 1 deterministic); 2,000-evaluation allowance.
Means include seeds with a feasible incumbent; feasibility counts are reported. Regret ≤ 10^{-12} is shown at 10^{-12}; early-stopped baseline incumbents are carried forward.