bbob mixint f17 d10 i1 — what the agent chose, and when

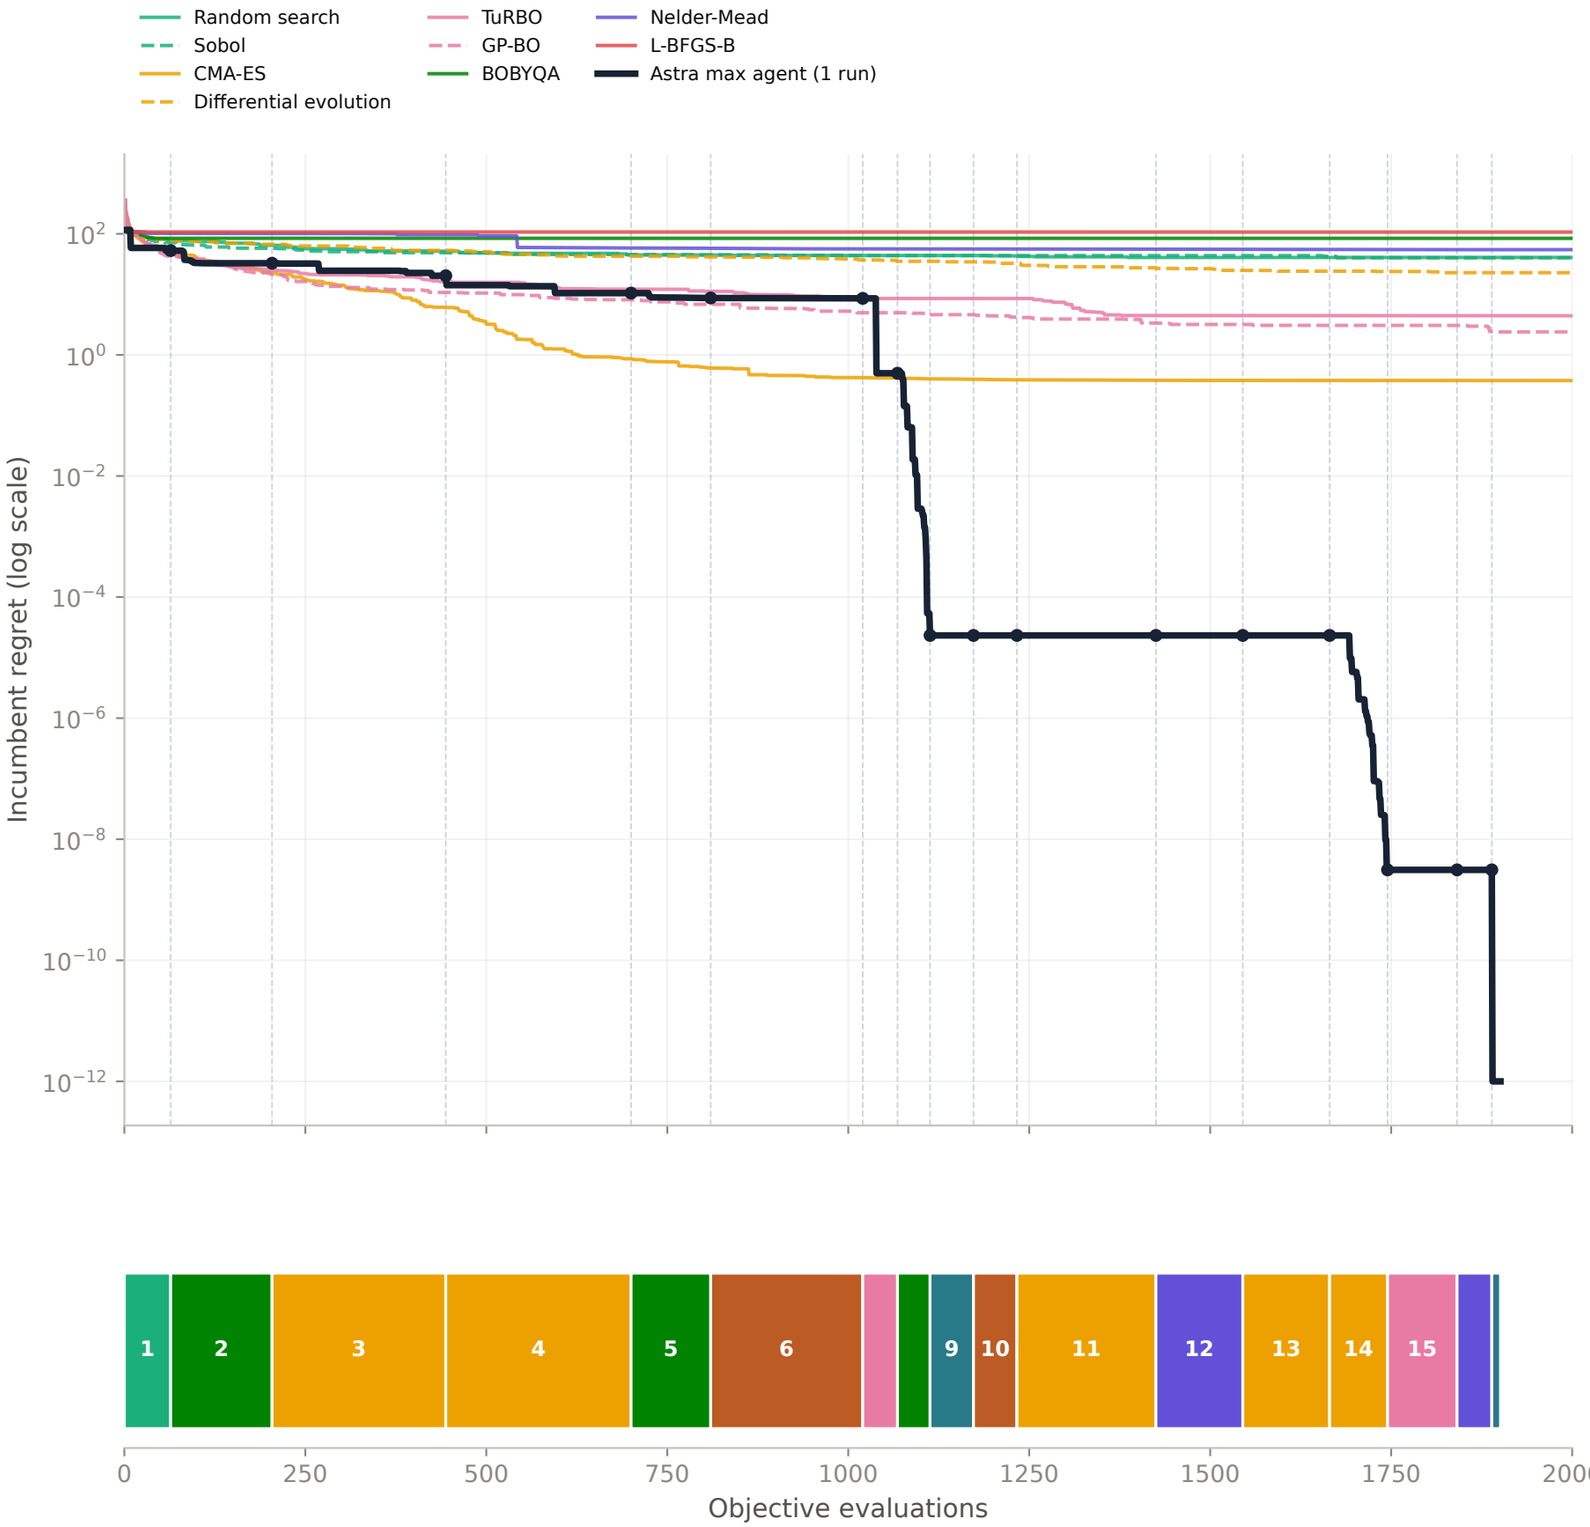

Block decisions

1. sobol_initial
   64 calls · evaluations 1–64
2. bobyqa_local1
   140 calls · evaluations 65–204
3. cma_global1
   240 calls · evaluations 205–444
4. cma_global1
   256 calls · evaluations 445–700
5. bobyqa_local2
   110 calls · evaluations 701–810
6. coordinate_scan
   210 calls · evaluations 811–1020
7. gp_two_coords
   48 calls · evaluations 1021–1068
8. bobyqa_two_coords
   45 calls · evaluations 1069–1113
9. plateau_coordinate_scan
   60 calls · evaluations 1114–1173
10. paired_plateau_moves
    60 calls · evaluations 1174–1233
11. cma_plateau_joint
    192 calls · evaluations 1234–1425
12. nelder_mead_pockets
    120 calls · evaluations 1426–1545
13. nelder_alternate_states
    120 calls · evaluations 1546–1665
14. nelder_precision
    80 calls · evaluations 1666–1745
15. gp_two_coords_rough
    96 calls · evaluations 1746–1841
16. surrogate_triple_moves
    48 calls · evaluations 1842–1889
17. precision_finish
    12 calls · evaluations 1890–1901

Boundaries: evaluator snapshots of prospective decisions; method changes inside a block stay in that block.
Agent: one fresh run, 1,901/2,000 evaluations; GPT-6 Astra, max reasoning. Baselines: seed means (10 stochastic / 1 deterministic); 2,000-evaluation allowance.
Regret ≤ 10^{-12} is shown at 10^{-12}; early-stopped baseline incumbents are carried forward.