bbob noisy f111 d5 i1 — what the agent chose, and when

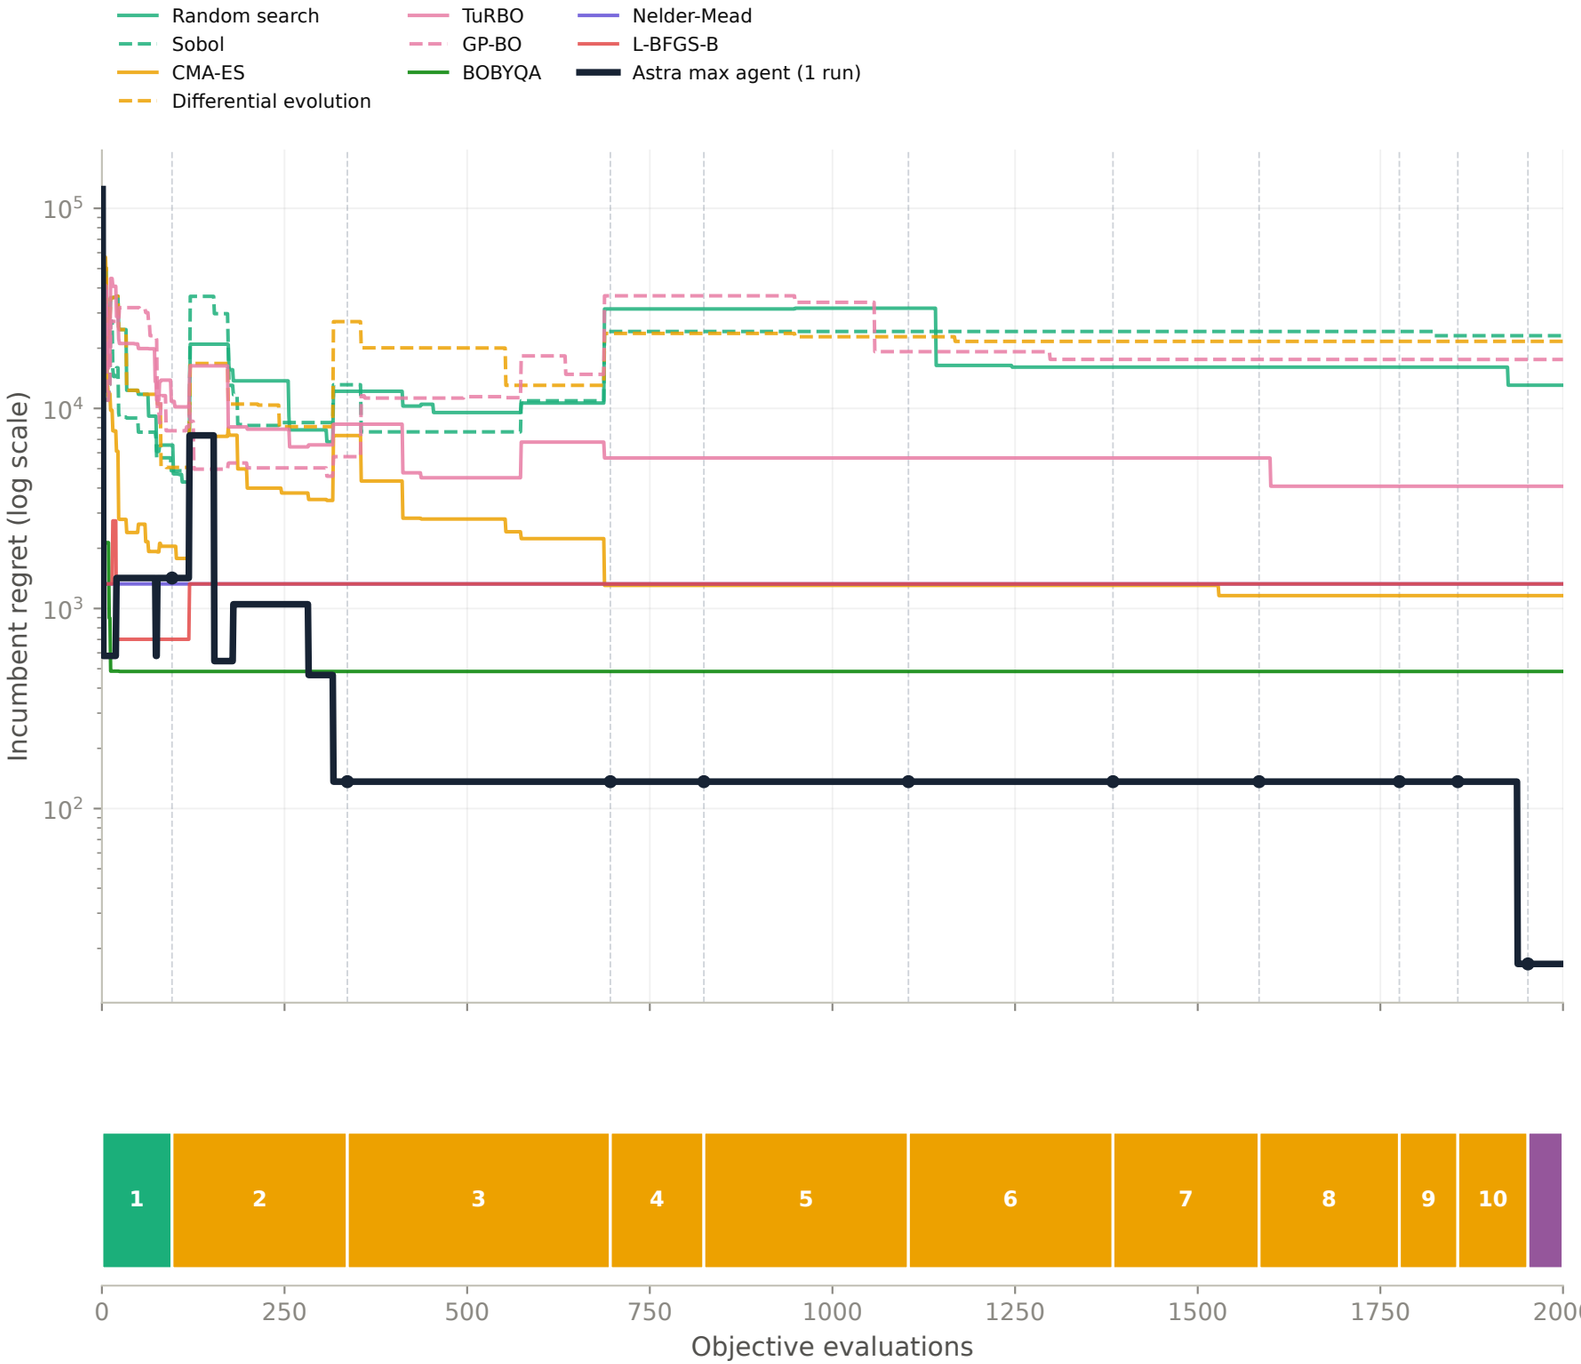

Block decisions

1. sobol_replicates
   96 calls · evaluations 1–96
2. cma_lower_quantile
   240 calls · evaluations 97–336
3. cma_lower_quantile
   360 calls · evaluations 337–696
4. quantile_model_probe
   128 calls · evaluations 697–824
5. cma_local_replicated
   280 calls · evaluations 825–1104
6. cma_local_replicated
   280 calls · evaluations 1105–1384
7. elite_model_replicates
   200 calls · evaluations 1385–1584
8. cma_elite_refine
   192 calls · evaluations 1585–1776
9. cma_elite_refine
   80 calls · evaluations 1777–1856
10. replicated_contenders
    96 calls · evaluations 1857–1952
11. local_interpolation
    48 calls · evaluations 1953–2000

Boundaries: evaluator snapshots of prospective decisions; method changes inside a block stay in that block.
Agent: one fresh run, 2,000/2,000 evaluations; GPT-6 Astra, max reasoning. Baselines: seed means (10 stochastic / 1 deterministic); 2,000-evaluation allowance.
Noisy task: score uses the noise-free value of the lowest-observed-y incumbent. Regret ≤ 10^{-12} is shown at 10^{-12}; early-stopped baseline incumbents are carried forward.