cec2017 f8 d30 — what the agent chose, and when

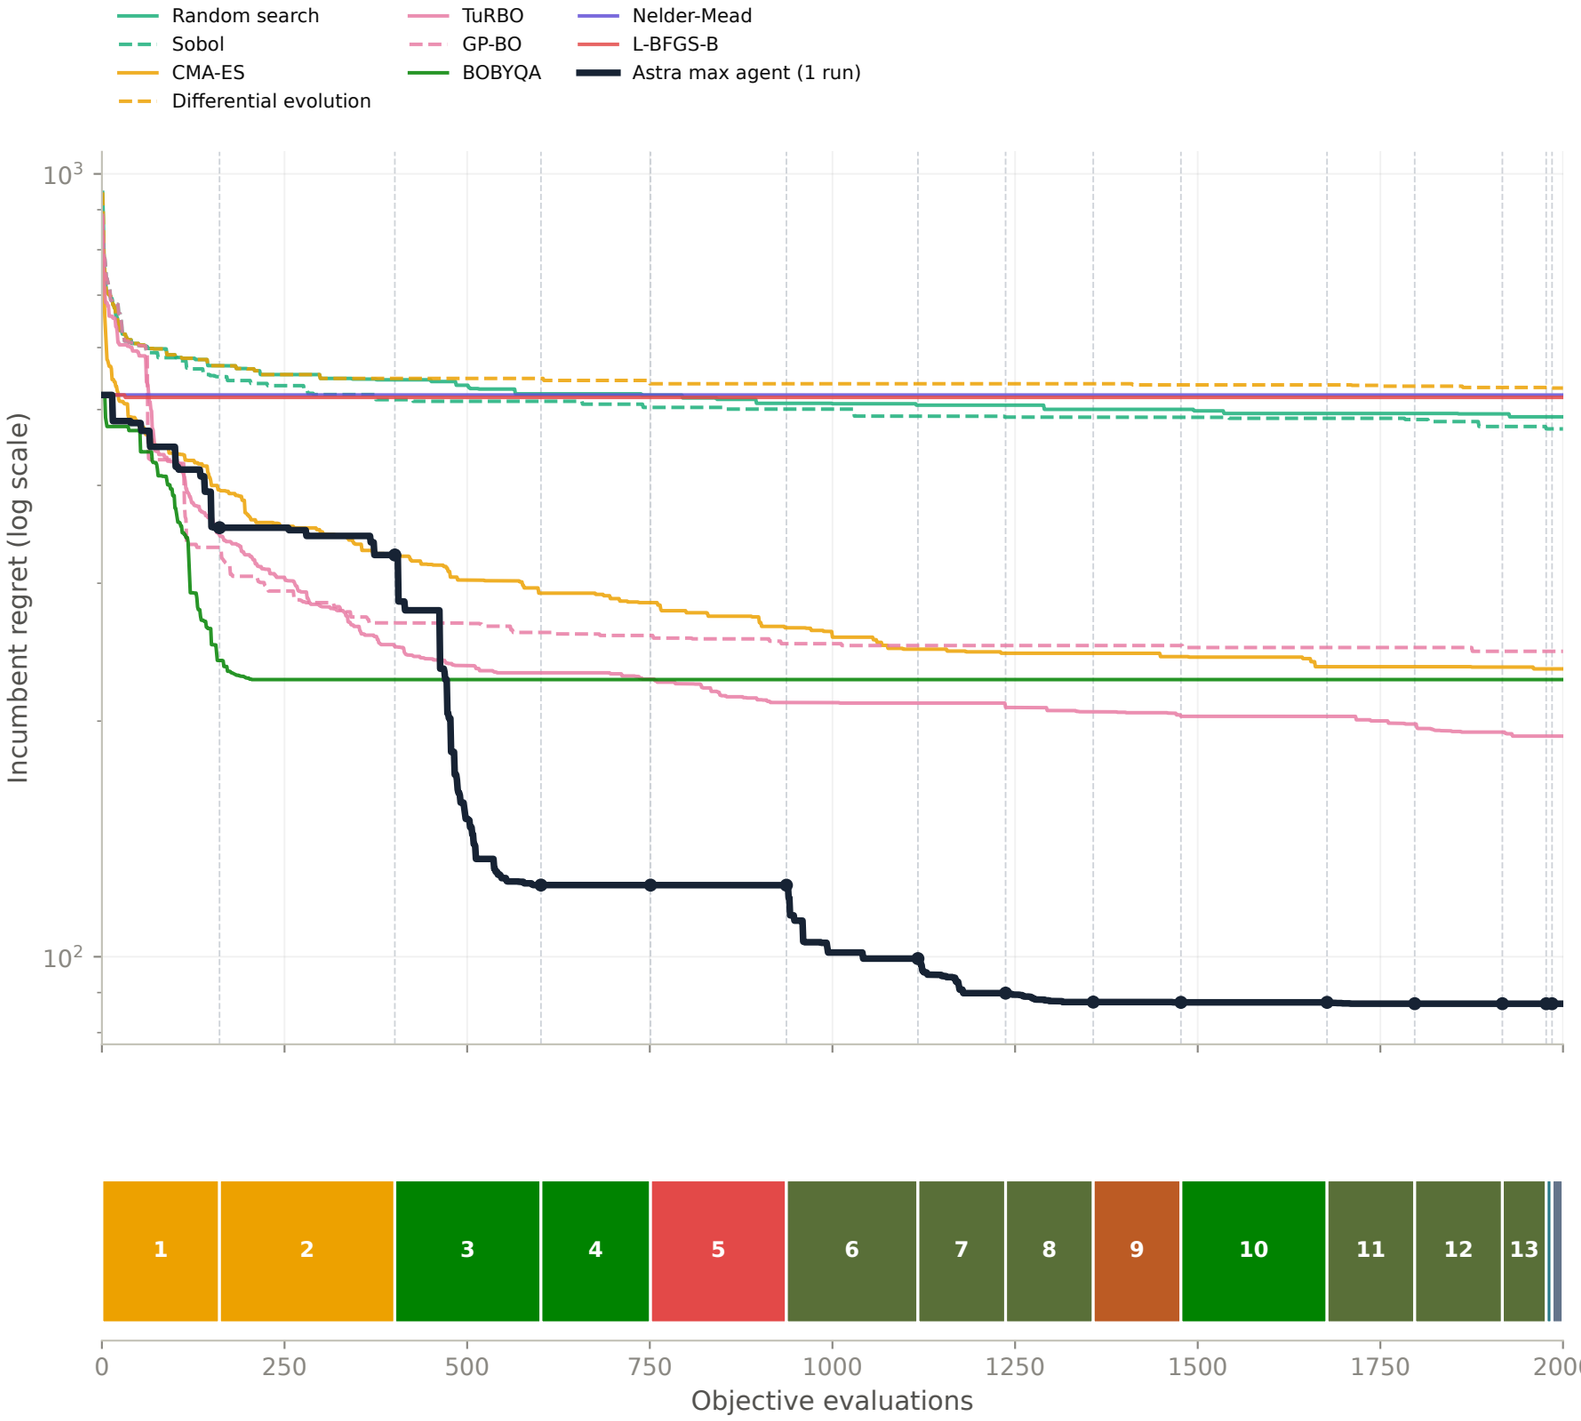

Block decisions

1. cma_es
   161 calls · evaluations 1–161
2. cma_es
   240 calls · evaluations 162–401
3. bobyqa_local
   200 calls · evaluations 402–601
4. bobyqa_local
   150 calls · evaluations 602–751
5. lbfgsb_local
   186 calls · evaluations 752–937
6. coordinate_poll
   180 calls · evaluations 938–1117
7. coordinate_poll
   120 calls · evaluations 1118–1237
8. coordinate_poll
   120 calls · evaluations 1238–1357
9. rotated_poll
   120 calls · evaluations 1358–1477
10. bobyqa_perturb
    200 calls · evaluations 1478–1677
11. coordinate_poll
    120 calls · evaluations 1678–1797
12. neutral_coord_poll
    120 calls · evaluations 1798–1917
13. neutral_coord_poll
    60 calls · evaluations 1918–1977
14. scalar_refine
    8 calls · evaluations 1978–1985
15. coordinate_grid
    15 calls · evaluations 1986–2000

Boundaries: evaluator snapshots of prospective decisions; method changes inside a block stay in that block.
Agent: one fresh run, 2,000/2,000 evaluations; GPT-6 Astra, max reasoning. Baselines: seed means (10 stochastic / 1 deterministic); 2,000-evaluation allowance.
Regret ≤ 10^{-12} is shown at 10^{-12}; early-stopped baseline incumbents are carried forward.