nevergrad noisy cigar d20 n0.1 — what the agent chose, and when

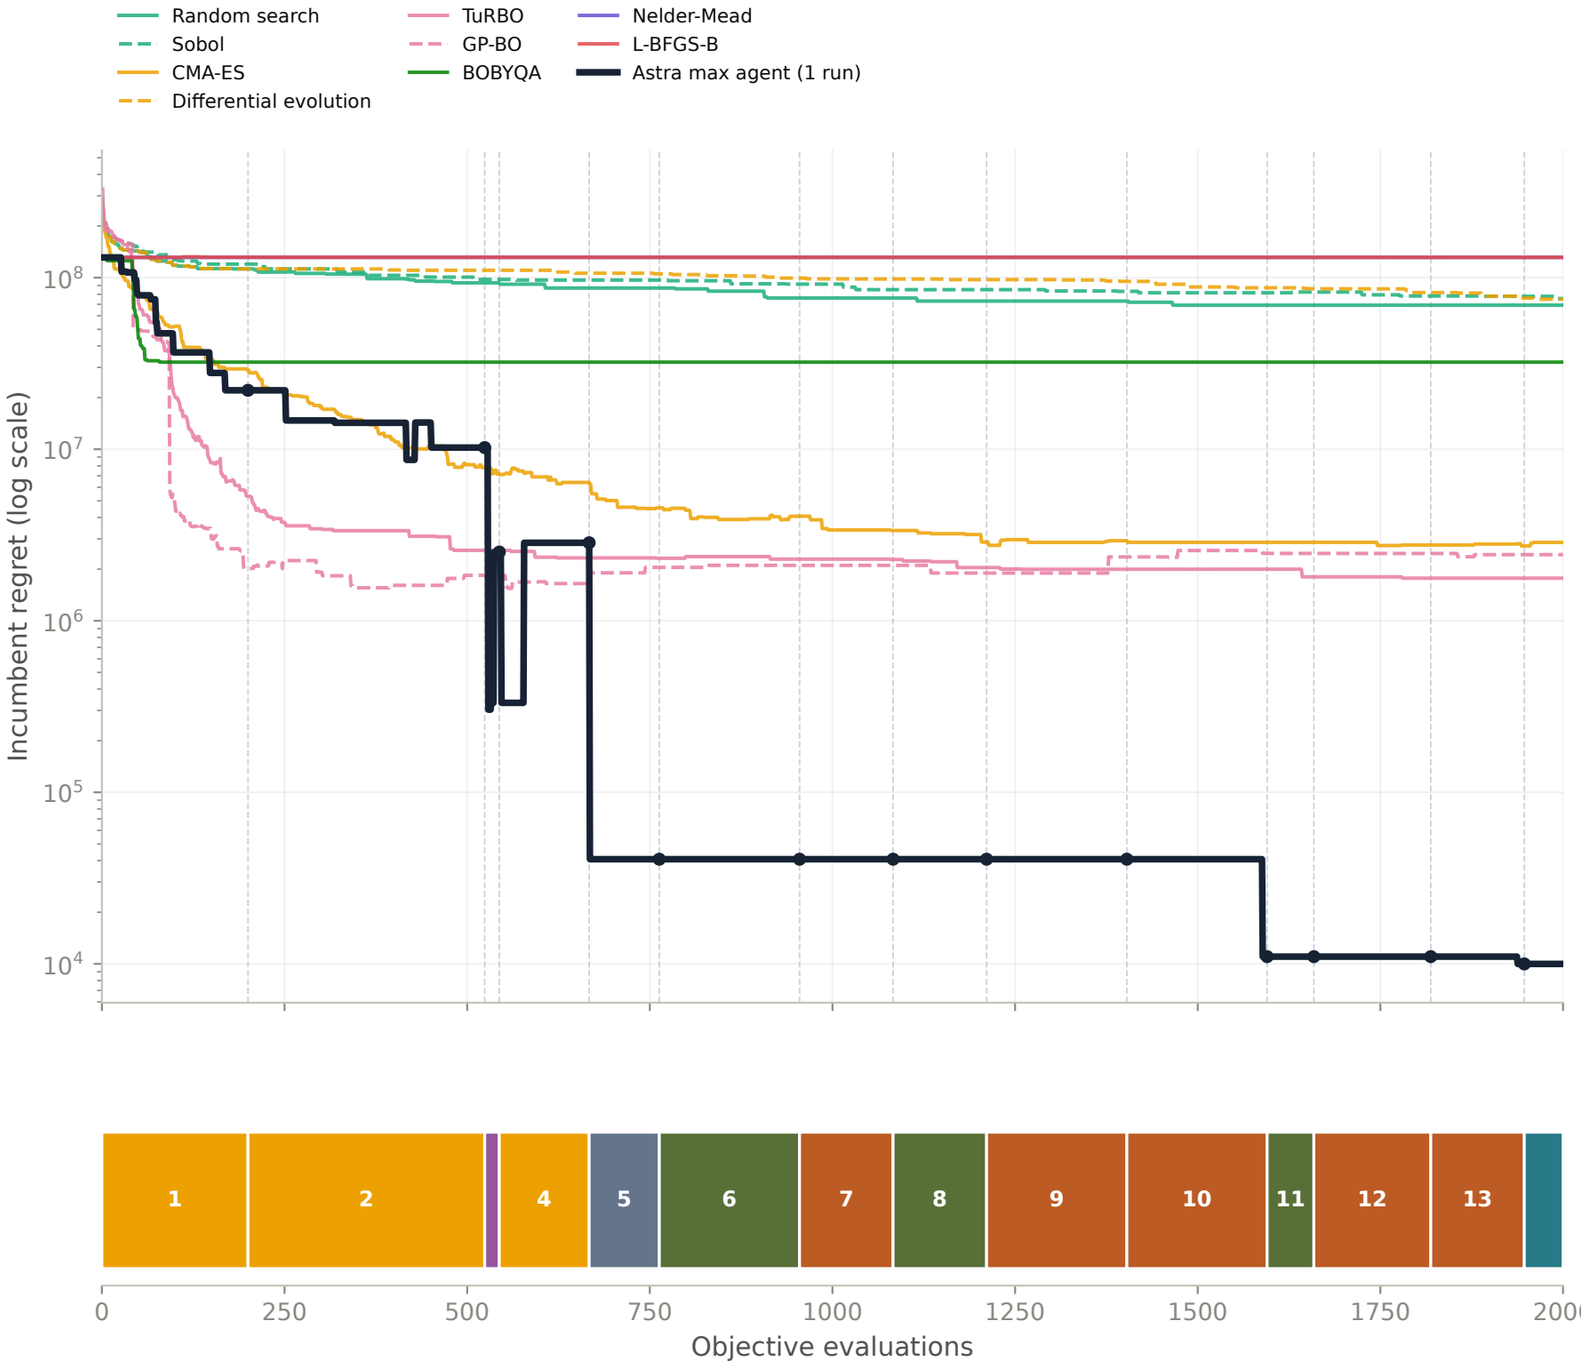

Block decisions

1. cma_noise_probe → cma_es
   200 calls · evaluations 1–200
2. cma_recheck → cma_es
   324 calls · evaluations 201–524
3. quadratic_probe
   20 calls · evaluations 525–544
4. quadratic_local_design
   123 calls · evaluations 545–667
5. diagonal_probe
   96 calls · evaluations 668–763
6. orthogonal_gradient
   192 calls · evaluations 764–955
7. diagonal_refine_replicate
   128 calls · evaluations 956–1083
8. orthogonal_gradient
   128 calls · evaluations 1084–1211
9. diagonal_refine_replicate
   192 calls · evaluations 1212–1403
10. diagonal_refine_replicate
    192 calls · evaluations 1404–1595
11. orthogonal_gradient
    64 calls · evaluations 1596–1659
12. diagonal_refine_replicate
    160 calls · evaluations 1660–1819
13. diagonal_refine_replicate
    128 calls · evaluations 1820–1947
14. incumbent_validation
    53 calls · evaluations 1948–2000

Boundaries: evaluator snapshots of prospective decisions; method changes inside a block stay in that block.
Agent: one fresh run, 2,000/2,000 evaluations; GPT-6 Astra, max reasoning. Baselines: seed means (10 stochastic / 1 deterministic); 2,000-evaluation allowance.
Noisy task: score uses the noise-free value of the lowest-observed-y incumbent. Regret ≤ 10^{-12} is shown at 10^{-12}; early-stopped baseline incumbents are carried forward.