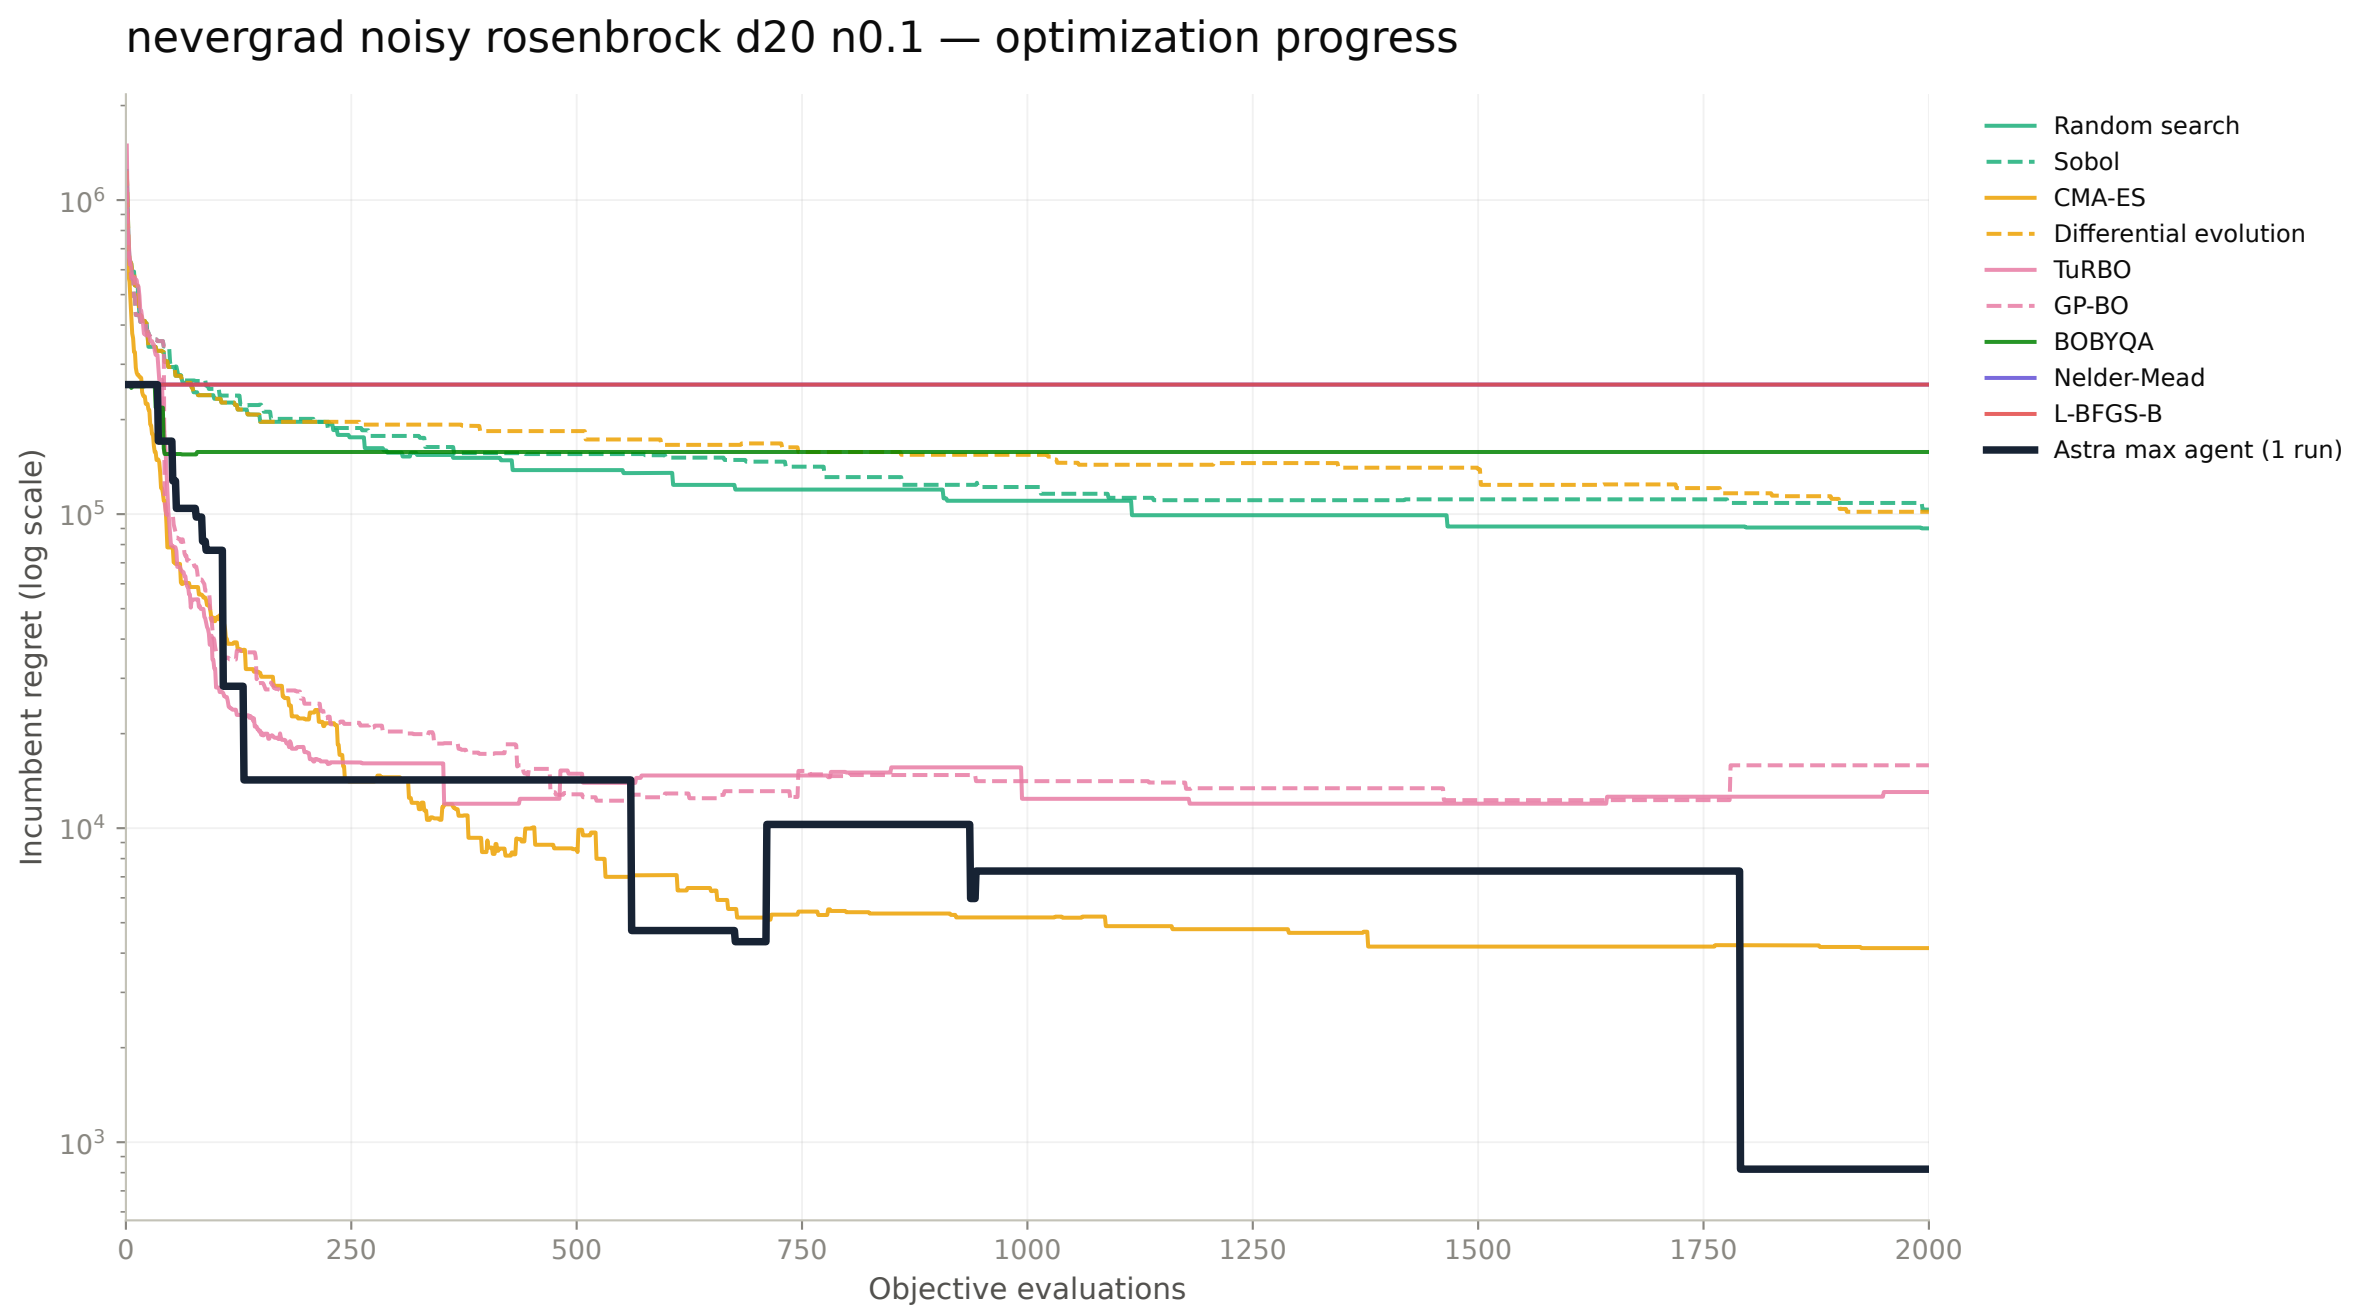

Agent: one fresh run, 2,000/2,000 evaluations; GPT-6 Astra, max reasoning. Baselines: seed means (10 stochastic / 1 deterministic); 2,000-evaluation allowance. Noisy task: score uses the noise-free value of the lowest-observed-y incumbent. Regret ≤ 10^{-12} is shown at 10^{-12}; early-stopped baseline incumbents are carried forward.