rastrigin d10 i1 — what the agent chose, and when

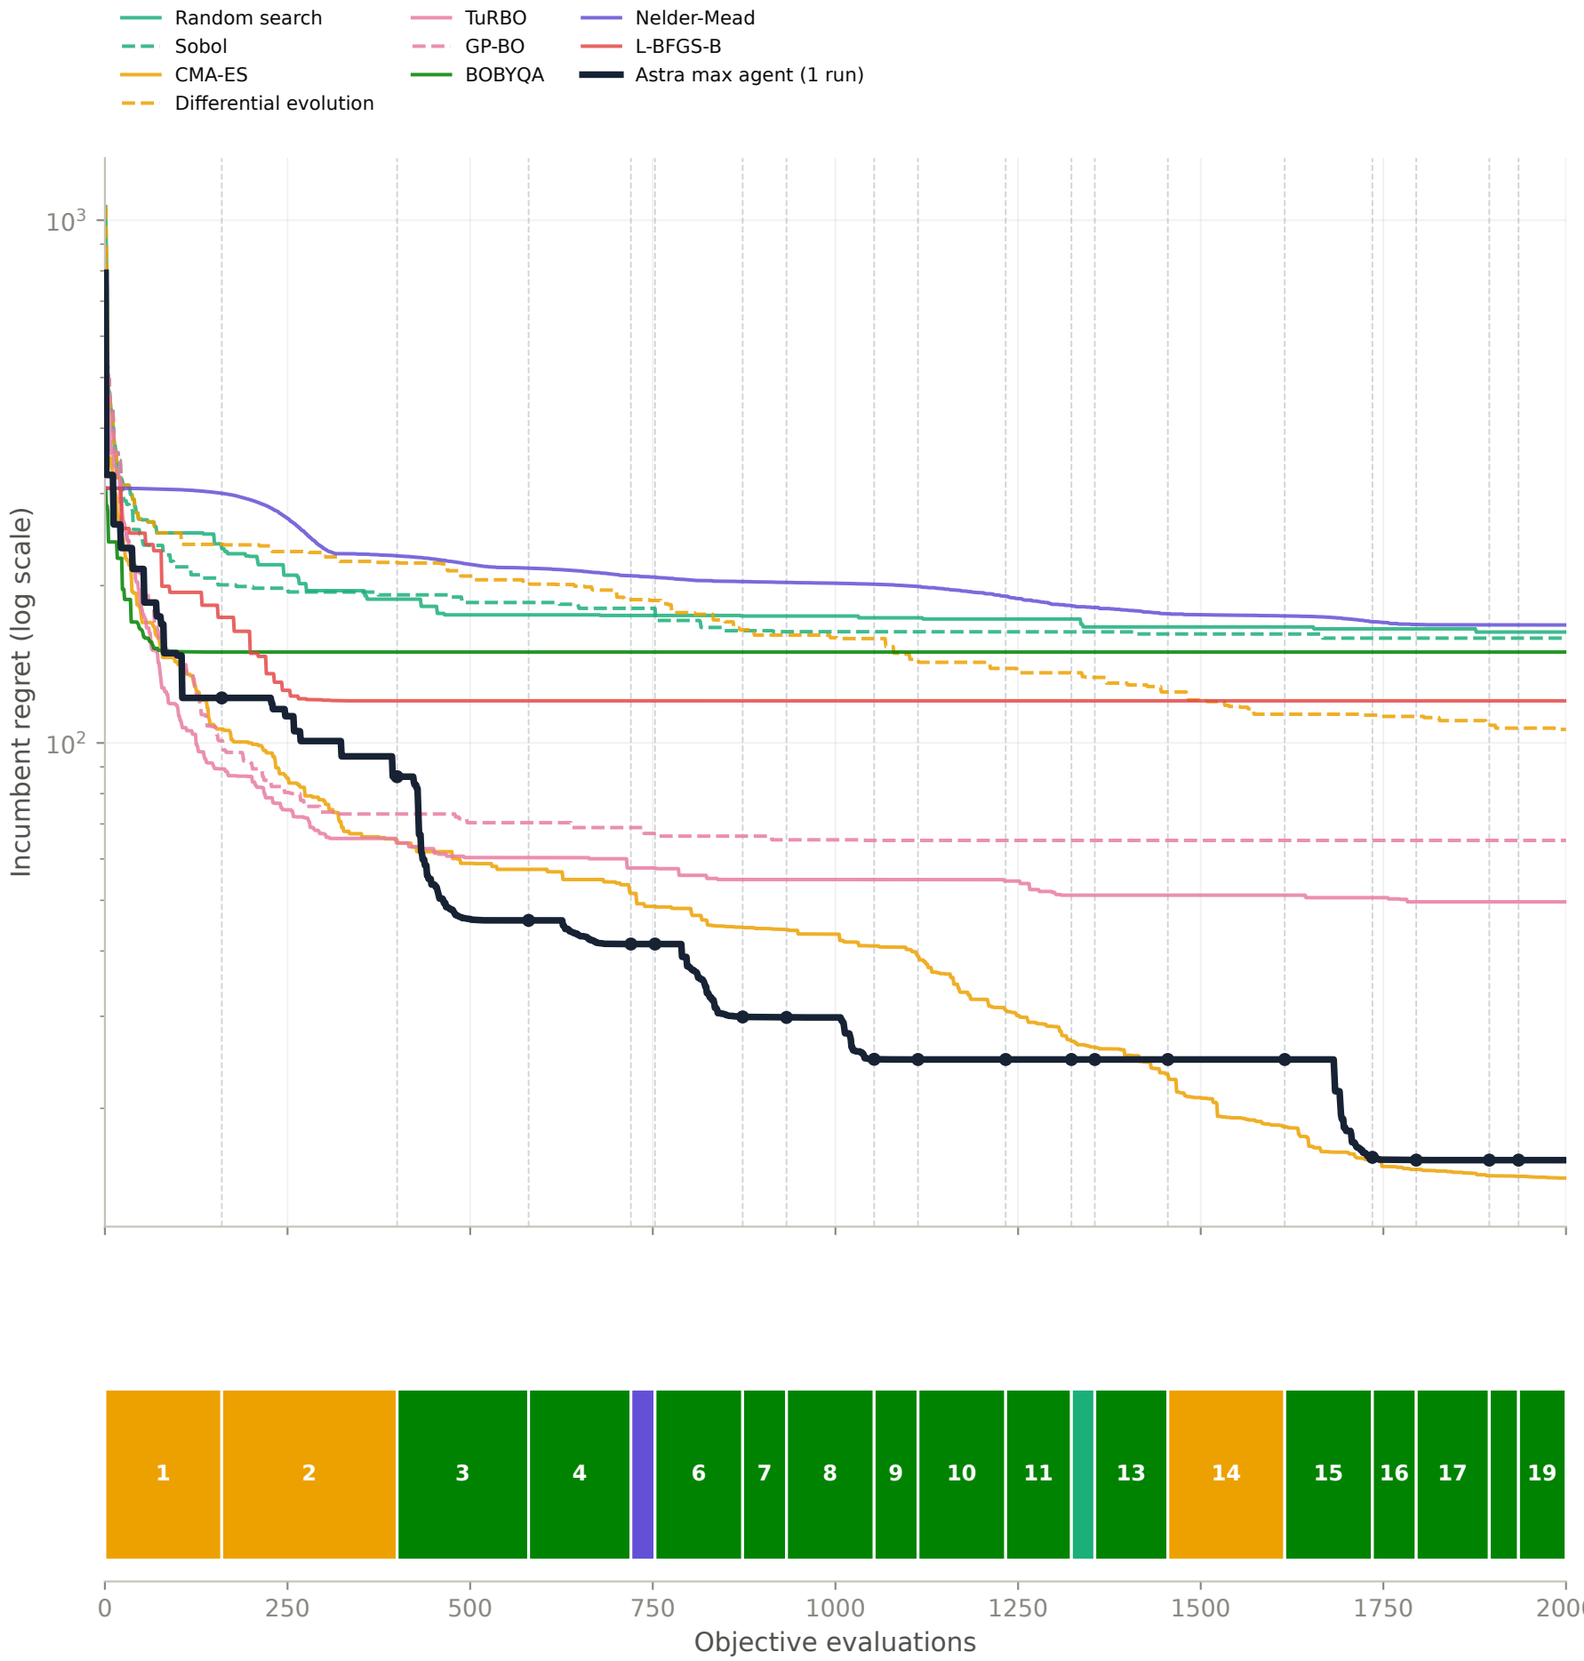

Block decisions

1. cma_es
   160 calls · evaluations 1–160
2. cma_es
   240 calls · evaluations 161–400
3. bobyqa_local
   180 calls · evaluations 401–580
4. bobyqa_local
   140 calls · evaluations 581–720
5. surrogate_sobol
   33 calls · evaluations 721–753
6. bobyqa_local
   120 calls · evaluations 754–873
7. bobyqa_local
   60 calls · evaluations 874–933
8. bobyqa_local
   120 calls · evaluations 934–1053
9. bobyqa_local
   60 calls · evaluations 1054–1113
10. bobyqa_local
    120 calls · evaluations 1114–1233
11. bobyqa_local
    90 calls · evaluations 1234–1323
12. incumbent_sobol
    32 calls · evaluations 1324–1355
13. bobyqa_local
    100 calls · evaluations 1356–1455
14. cma_es
    160 calls · evaluations 1456–1615
15. bobyqa_local
    120 calls · evaluations 1616–1735
16. bobyqa_local
    60 calls · evaluations 1736–1795
17. bobyqa_local
    100 calls · evaluations 1796–1895
18. bobyqa_local
    40 calls · evaluations 1896–1935
19. bobyqa_local
    65 calls · evaluations 1936–2000

Boundaries: evaluator snapshots of prospective decisions; method changes inside a block stay in that block.
Agent: one fresh run, 2,000/2,000 evaluations; GPT-6 Astra, max reasoning. Baselines: seed means (10 stochastic / 1 deterministic); 2,000-evaluation allowance.
Regret ≤ 10^{-12} is shown at 10^{-12}; early-stopped baseline incumbents are carried forward.