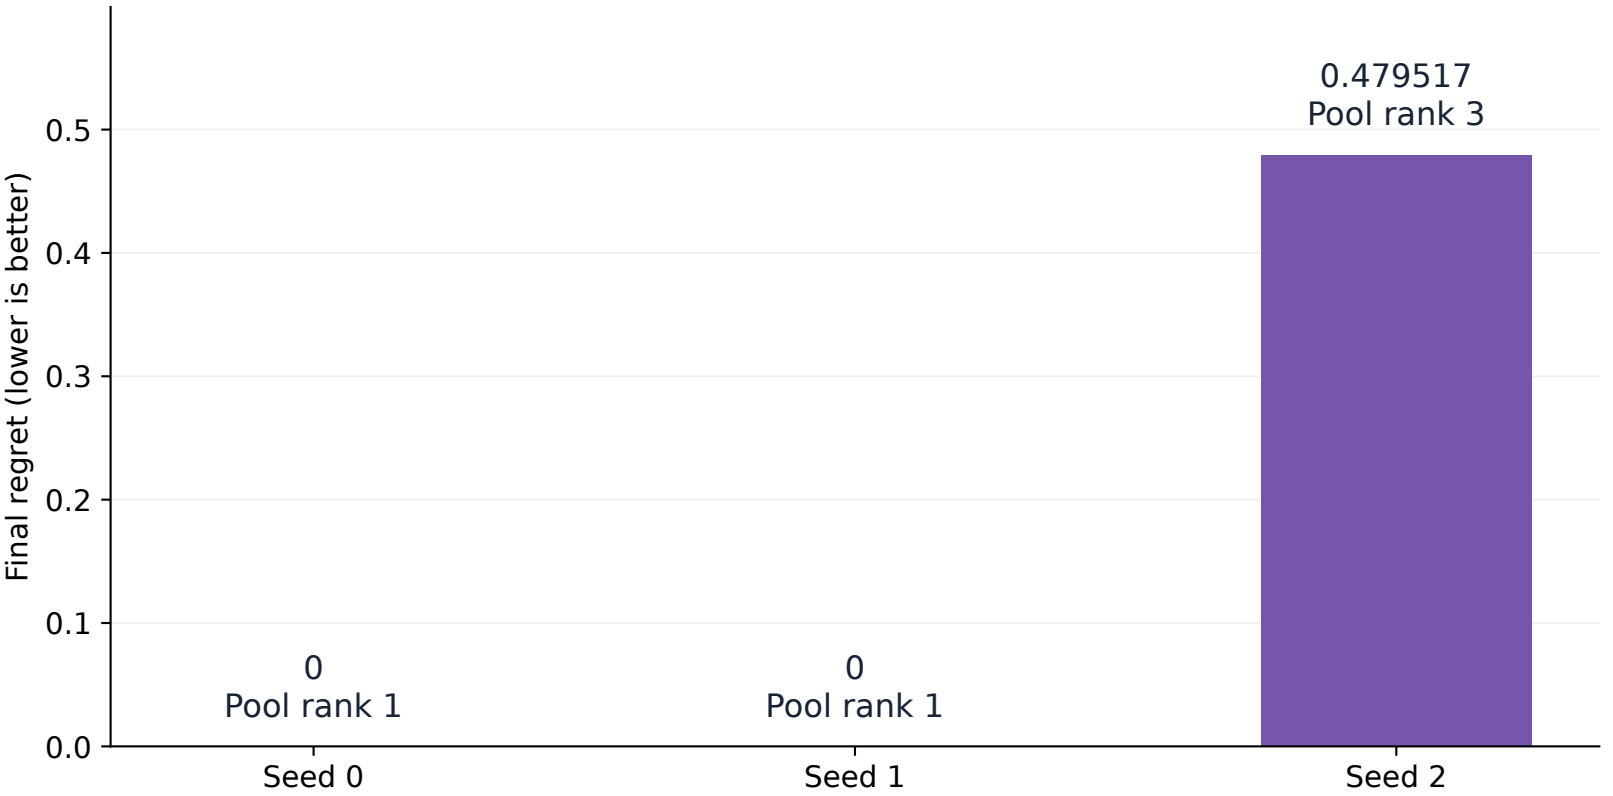

GPT-6 Astra · max reasoning · 100 evaluations per seed, including initialization. Seeds 0, 1, 2 are independent campaigns; model generation is not deterministic. Agent seed traces; standalone baselines appear in the complete comparison figure. Regret = pool maximum − best measured score.