MOF candidate pool — what each agent chose, and when

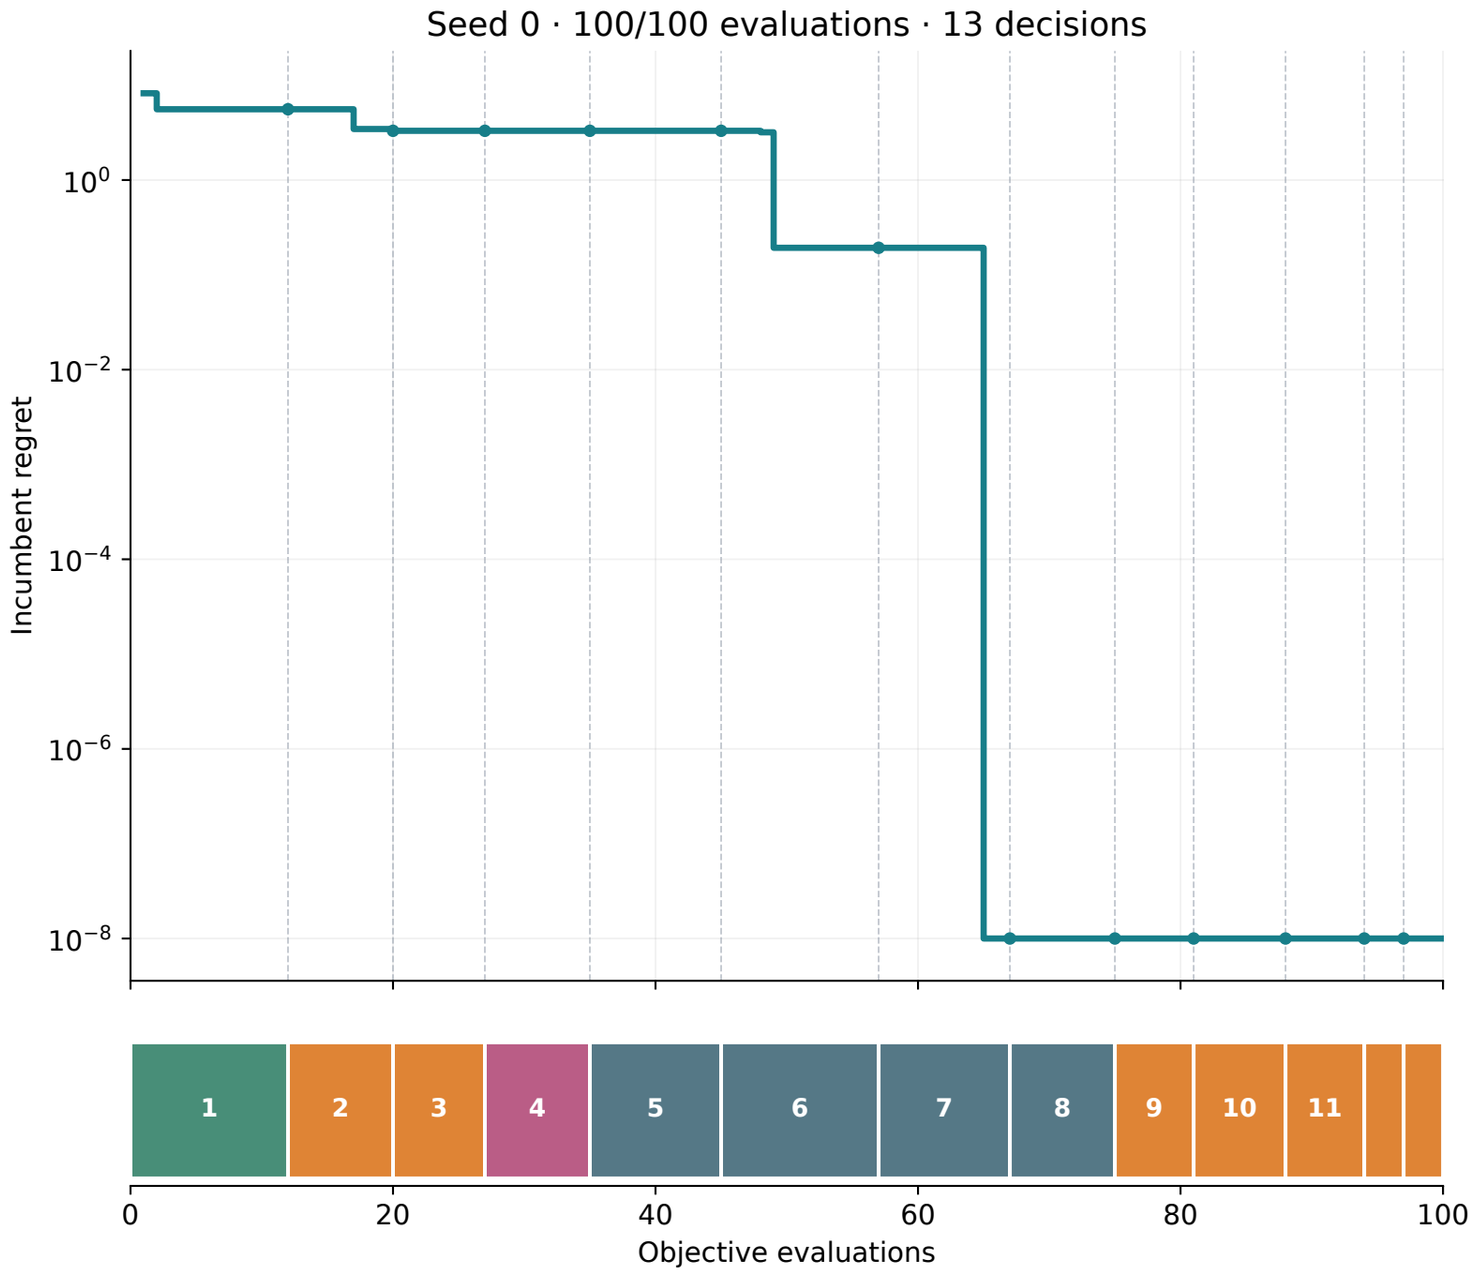

Block decisions

1. balanced pool design
   12 calls · evaluations 1–12
2. pool gp logei
   8 calls · evaluations 13–20
3. pool gp logei
   7 calls · evaluations 21–27
4. pool matern logei
   8 calls · evaluations 28–35
5. pool extratrees lcb
   10 calls · evaluations 36–45
6. pool extratrees lcb
   12 calls · evaluations 46–57
7. pool extratrees lcb
   10 calls · evaluations 58–67
8. pool extratrees lcb
   8 calls · evaluations 68–75
9. pool gp logei
   6 calls · evaluations 76–81
10. pool gp lcb
    7 calls · evaluations 82–88
11. pool gp tree lcb
    6 calls · evaluations 89–94
12. pool gp tree lcb
    3 calls · evaluations 95–97
13. pool gp tree lcb
    3 calls · evaluations 98–100

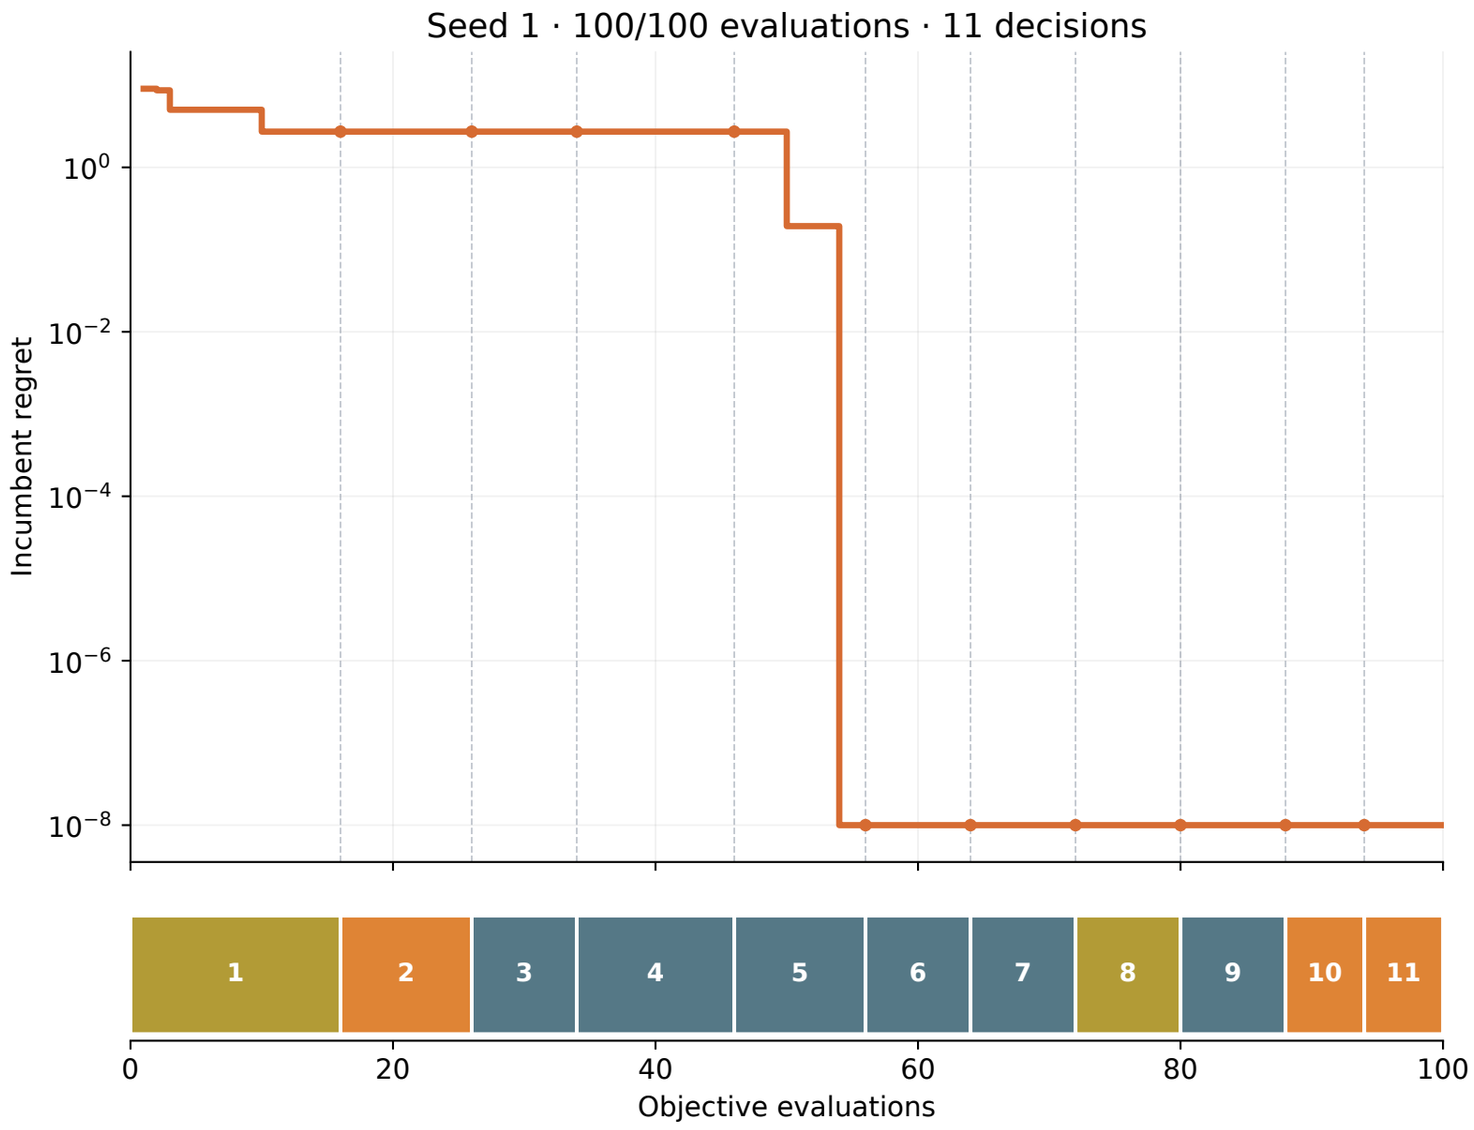

Block decisions

1. diverse initial
   16 calls · evaluations 1–16
2. pool gp logei
   10 calls · evaluations 17–26
3. extra trees lcb
   8 calls · evaluations 27–34
4. extra trees lcb
   12 calls · evaluations 35–46
5. extra trees lcb
   10 calls · evaluations 47–56
6. extra trees lcb
   8 calls · evaluations 57–64
7. extra trees lcb
   8 calls · evaluations 65–72
8. extra trees tail lcb
   8 calls · evaluations 73–80
9. extra trees lcb
   8 calls · evaluations 81–88
10. gp forest committee
    6 calls · evaluations 89–94
11. gp forest committee
    6 calls · evaluations 95–100

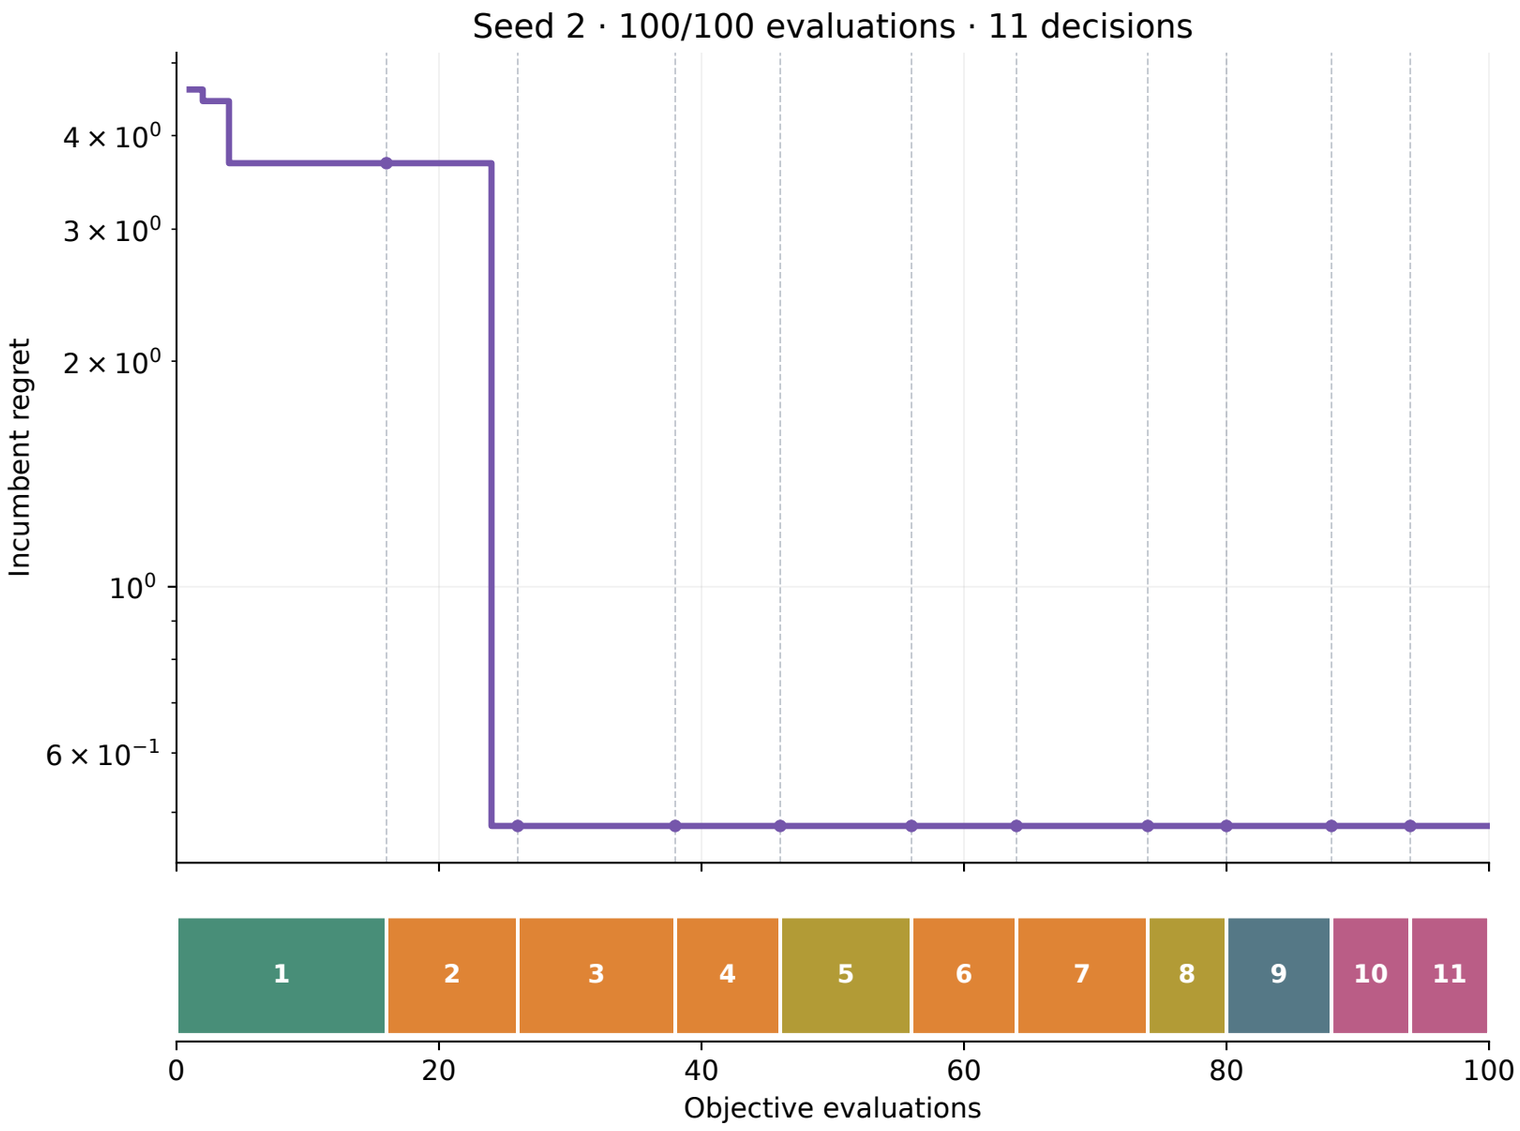

Block decisions

1. cluster design
   16 calls · evaluations 1–16
2. pool gp ei
   10 calls · evaluations 17–26
3. pool gp ei
   12 calls · evaluations 27–38
4. pool gp lcb
   8 calls · evaluations 39–46
5. forest lcb
   10 calls · evaluations 47–56
6. sparse gp lcb
   8 calls · evaluations 57–64
7. sparse gp lcb
   10 calls · evaluations 65–74
8. nearest incumbent
   6 calls · evaluations 75–80
9. sparse ensemble ei
   8 calls · evaluations 81–88
10. sparse ensemble lcb
    6 calls · evaluations 89–94
11. sparse ensemble lcb
    6 calls · evaluations 95–100

GPT-6 Astra · max reasoning · 100 evaluations per seed, including initialization. Seeds 0, 1, 2 are independent campaigns; model generation is not deterministic.
Agent seed traces; standalone baselines appear in the complete comparison figure. Regret = pool maximum − best measured score.
Boundaries come from prospective decision snapshots, including continuations. Zero regret is drawn at 1e-08.