MOF candidate pool — seed 1: decisions and progress

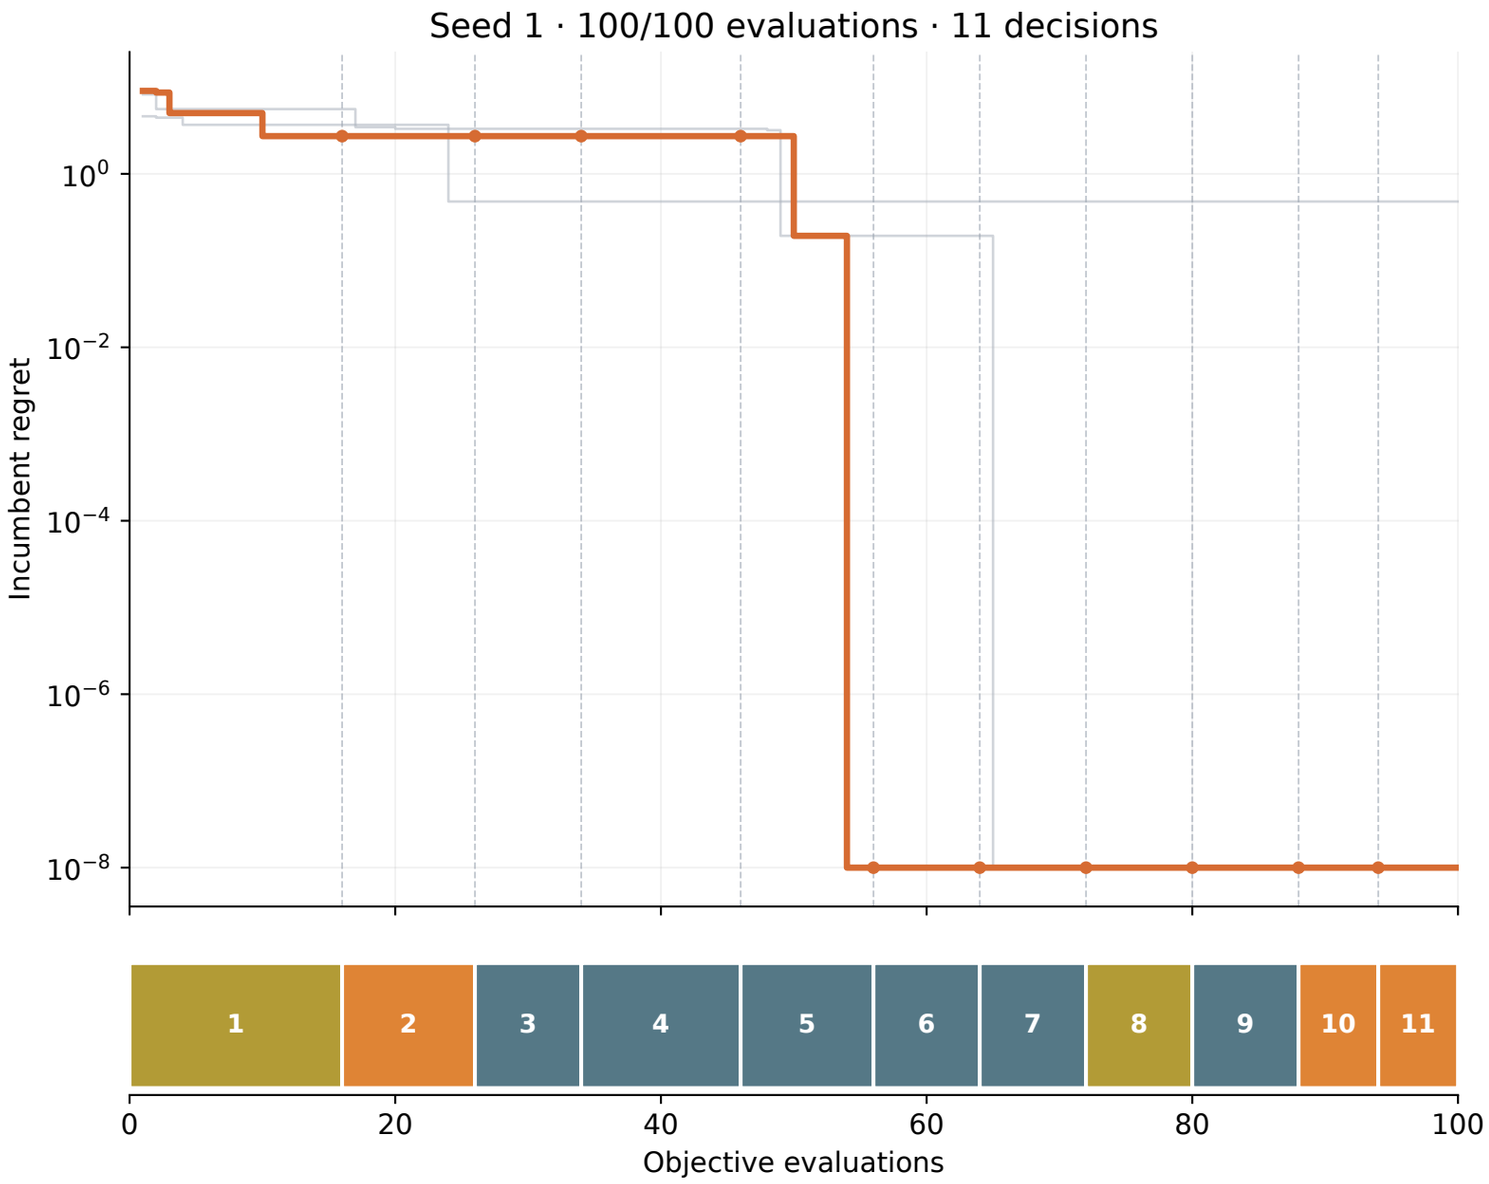

Block decisions

1. diverse initial
   16 calls · evaluations 1–16
2. pool gp logei
   10 calls · evaluations 17–26
3. extra trees lcb
   8 calls · evaluations 27–34
4. extra trees lcb
   12 calls · evaluations 35–46
5. extra trees lcb
   10 calls · evaluations 47–56
6. extra trees lcb
   8 calls · evaluations 57–64
7. extra trees lcb
   8 calls · evaluations 65–72
8. extra trees tail lcb
   8 calls · evaluations 73–80
9. extra trees lcb
   8 calls · evaluations 81–88
10. gp forest committee
    6 calls · evaluations 89–94
11. gp forest committee
    6 calls · evaluations 95–100

GPT-6 Astra · max reasoning · 100 evaluations per seed, including initialization. Seeds 0, 1, 2 are independent campaigns; model generation is not deterministic.
Agent seed traces; standalone baselines appear in the complete comparison figure. Regret = pool maximum − best measured score.
Boundaries come from prospective decision snapshots; gray curves are the other two agent seeds. Zero regret is drawn at 1e-08.