OPV device pool — what each agent chose, and when

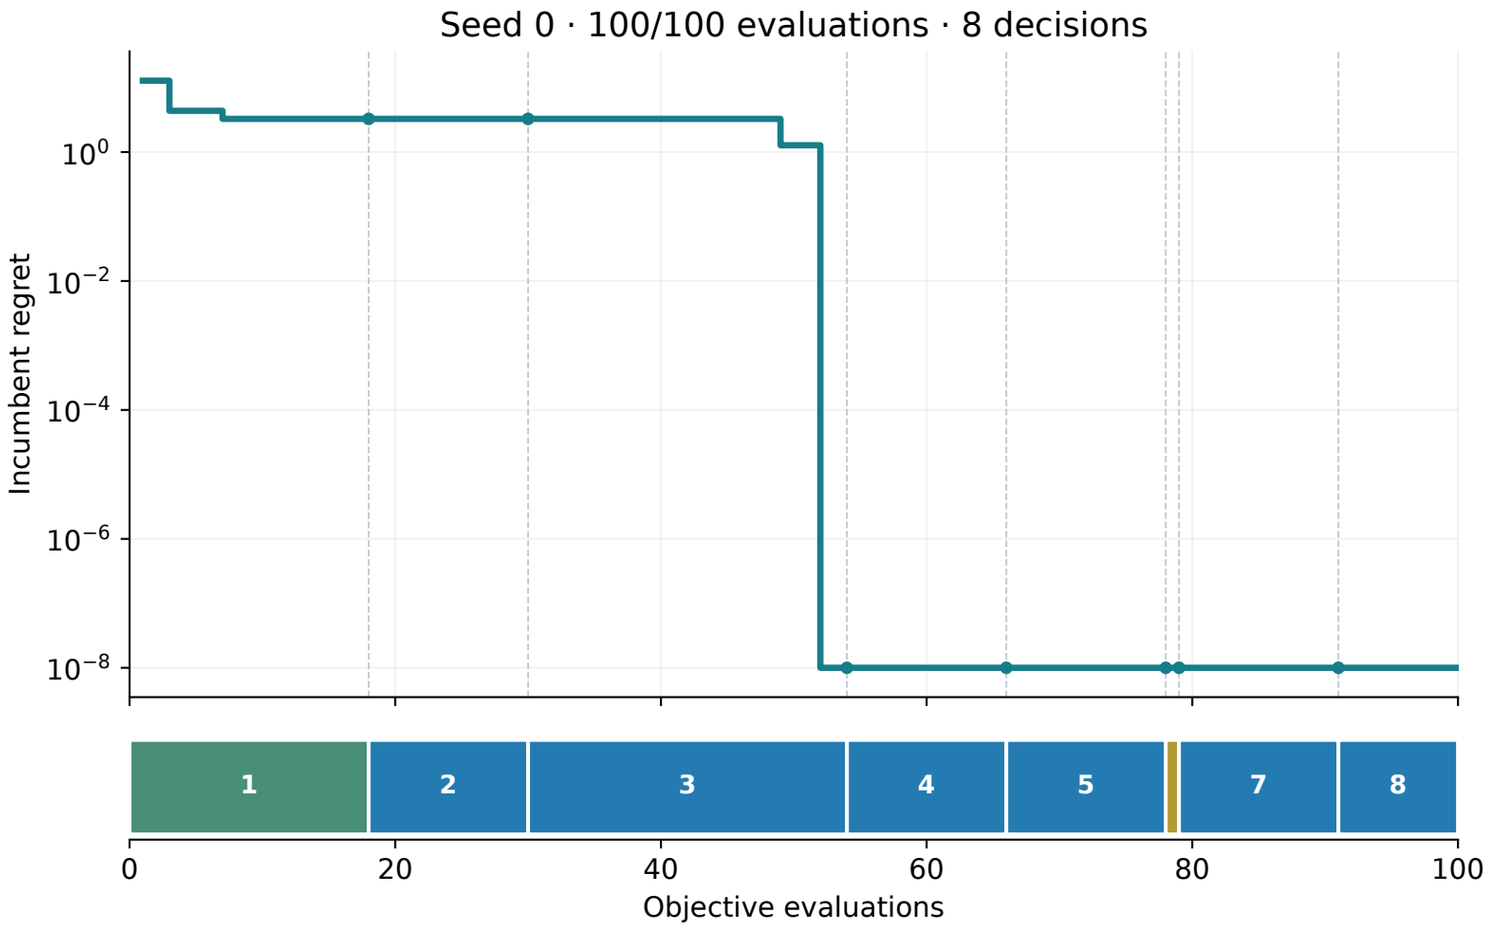

Block decisions

1. sobol reference
   18 calls · evaluations 1–18
2. cma es
   12 calls · evaluations 19–30
3. cma es
   24 calls · evaluations 31–54
4. cma es
   12 calls · evaluations 55–66
5. cma es
   12 calls · evaluations 67–78
6. local reference
   1 calls · evaluations 79–79
7. cma es
   12 calls · evaluations 80–91
8. cma final design
   9 calls · evaluations 92–100

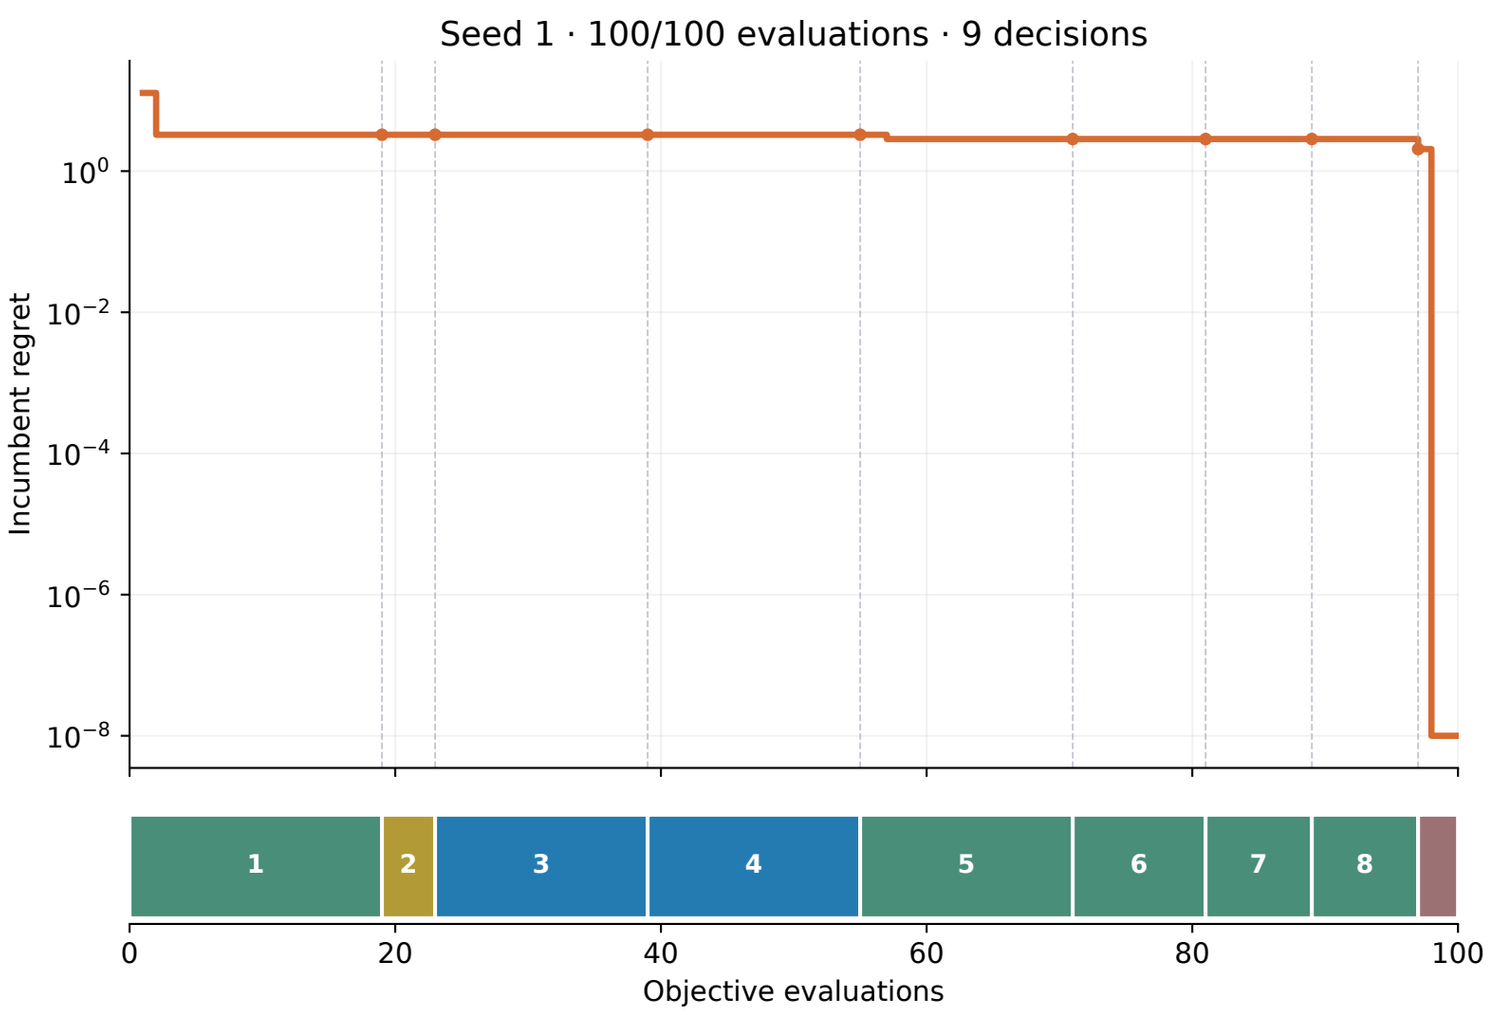

Block decisions

1. sobol anchors
   19 calls · evaluations 1–19
2. paired recheck
   4 calls · evaluations 20–23
3. diag cma
   16 calls · evaluations 24–39
4. diag cma
   16 calls · evaluations 40–55
5. sobol global
   16 calls · evaluations 56–71
6. local sobol recheck
   10 calls · evaluations 72–81
7. sobol global
   8 calls · evaluations 82–89
8. sobol global
   8 calls · evaluations 90–97
9. record local probe
   3 calls · evaluations 98–100

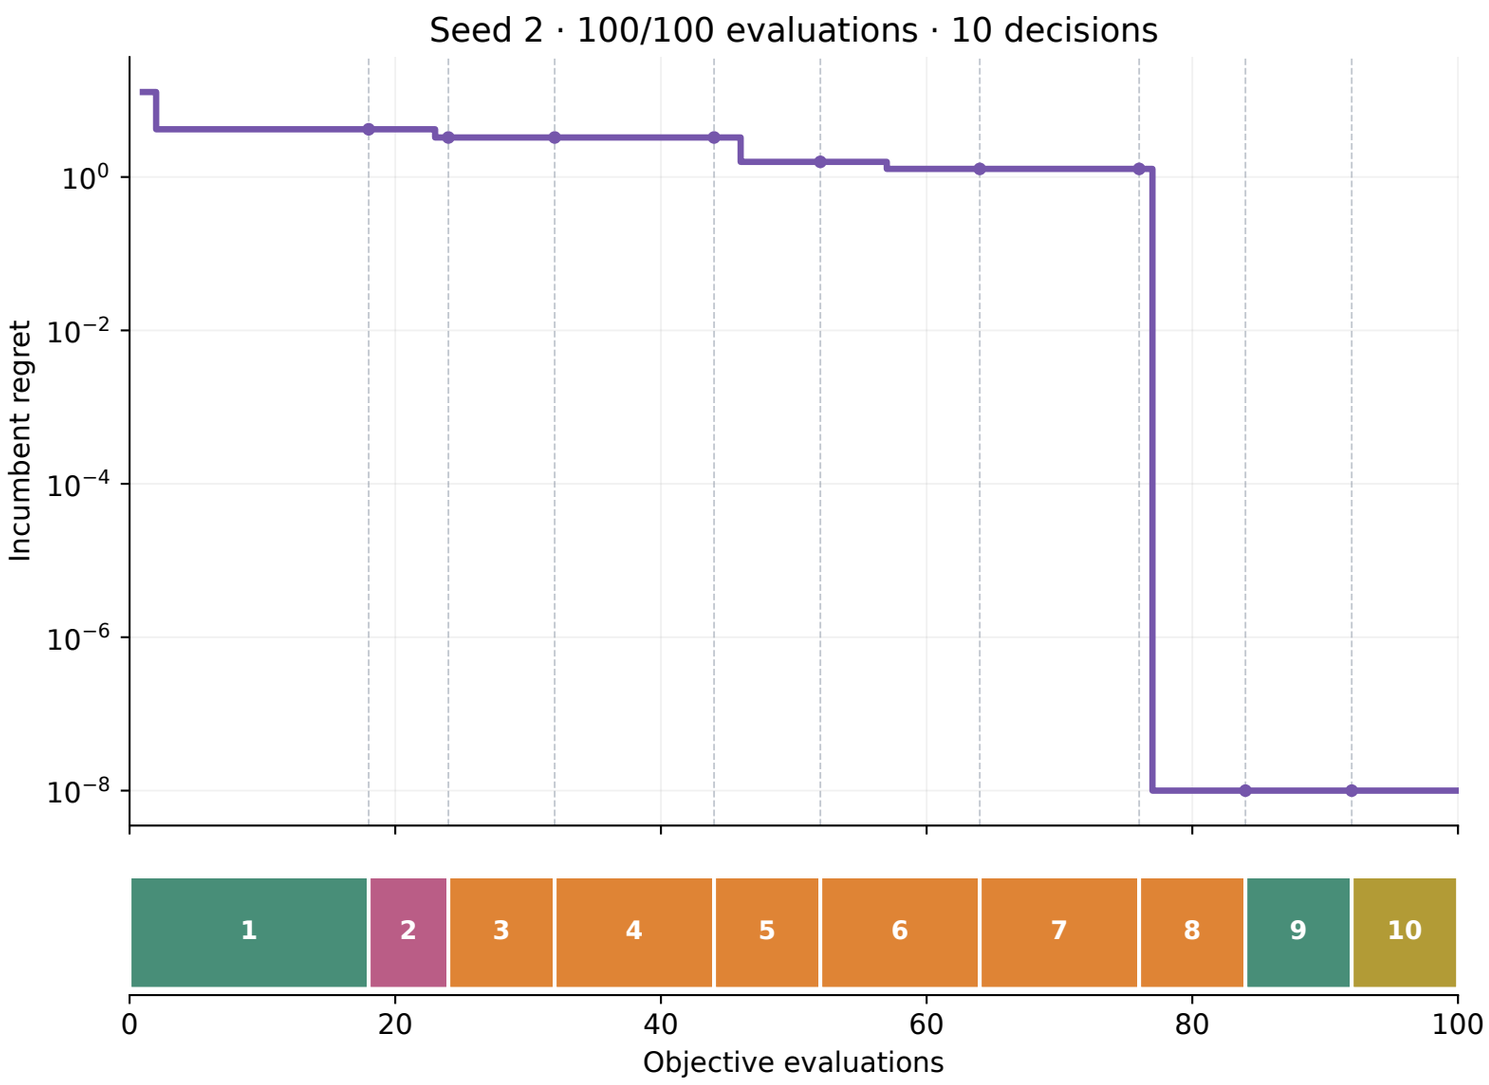

Block decisions

1. sobol anchor
   18 calls · evaluations 1–18
2. paired history
   6 calls · evaluations 19–24
3. gp history lcb
   8 calls · evaluations 25–32
4. gp history lcb
   12 calls · evaluations 33–44
5. gp history exploit + history control
   8 calls · evaluations 45–52
6. gp history exploit + local sobol
   12 calls · evaluations 53–64
7. gp history exploit + local sobol
   12 calls · evaluations 65–76
8. incumbent revisit + gp history exploit
   8 calls · evaluations 77–84
9. face sobol + incumbent revisit
   8 calls · evaluations 85–92
10. elite revisits
    8 calls · evaluations 93–100

GPT-6 Astra · max reasoning · 100 evaluations per seed, including initialization. Seeds 0, 1, 2 are independent campaigns; model generation is not deterministic.
Agent seed traces; standalone baselines appear in the complete comparison figure. Regret = pool maximum − best measured score.
Boundaries come from prospective decision snapshots, including continuations. Zero regret is drawn at 1e-08.