OPV device pool — seed 2: decisions and progress

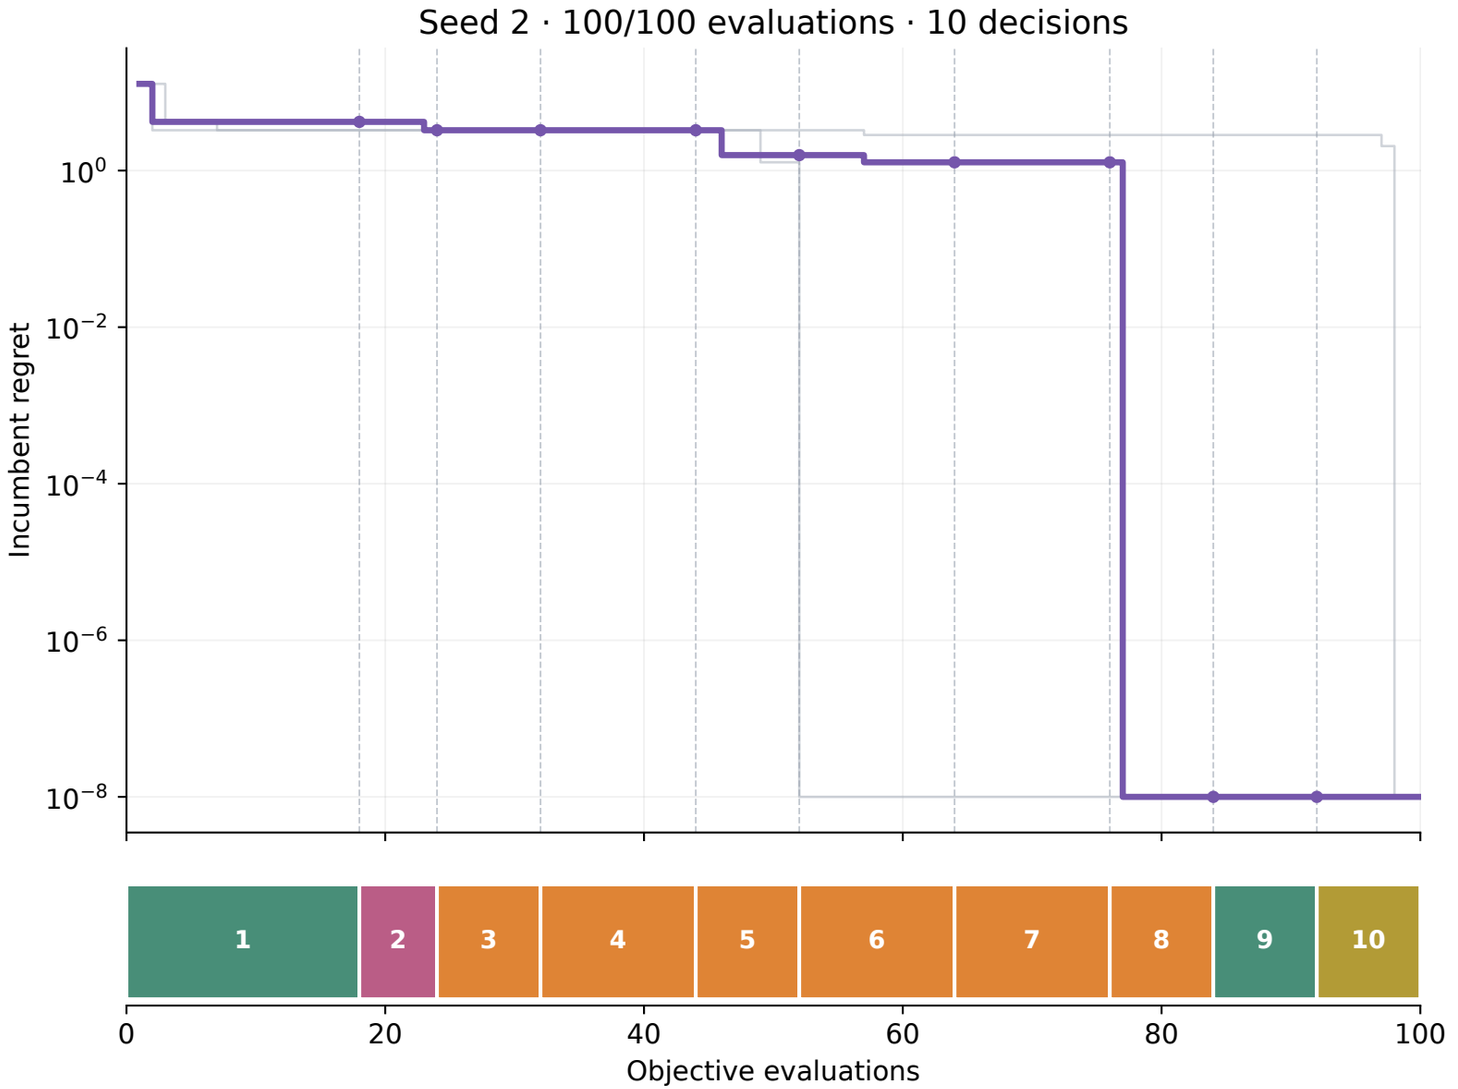

Block decisions

1. sobol anchor
   18 calls · evaluations 1–18

2. paired history
   6 calls · evaluations 19–24

3. gp history lcb
   8 calls · evaluations 25–32

4. gp history lcb
   12 calls · evaluations 33–44

5. gp history exploit + history control
   8 calls · evaluations 45–52

6. gp history exploit + local sobol
   12 calls · evaluations 53–64

7. gp history exploit + local sobol
   12 calls · evaluations 65–76

8. incumbent revisit + gp history exploit
   8 calls · evaluations 77–84

9. face sobol + incumbent revisit
   8 calls · evaluations 85–92

10. elite revisits
    8 calls · evaluations 93–100

GPT-6 Astra · max reasoning · 100 evaluations per seed, including initialization. Seeds 0, 1, 2 are independent campaigns; model generation is not deterministic.
Agent seed traces; standalone baselines appear in the complete comparison figure. Regret = pool maximum − best measured score.
Boundaries come from prospective decision snapshots; gray curves are the other two agent seeds. Zero regret is drawn at 1e-08.